

PARTNER ARCHITECTURE BRIEFING · PROPOSED INTEGRATION

Ugence + ServiceNow

Independent Action-Level Authority for Governed Agentic Workflows

Detailed Architecture and Use-Case Briefing

Extending ServiceNow-governed workflows with independently verifiable authorization, operational clearance, execution reconciliation, and governed-value evidence.

Rakesh Mohan

Founder, Ugence Labs · ugence.ai · August 2026

All ServiceNow integrations described are proposed. No ServiceNow connector currently ships. The scenarios are illustrative enterprise workflows, not claimed Ugence customer deployments.

Contents

1. Executive summary	3
2. What ServiceNow already provides	5
3. Proposed ServiceNow–Ugence responsibility boundary	6
4. Understanding the data journey	7
5. Ugence module responsibilities	10
6. Detailed lead scenario: UC-5 autonomous change execution	14
7. Additional detailed enterprise scenarios	20
7.1 UC-11 — Vulnerability remediation and emergency patching	20
7.2 UC-6 — Access provisioning with segregation of duties	22
7.3 UC-3 — High-risk AI action enforcement	24
7.4 UC-4 — External agents through Action Fabric, A2A and MCP	26
8. Full use-case portfolio	28
9. Honest overlap and precise differentiation	31
10. Enterprise Governed Value	32
11. Proposed bounded pilot	33
12. Discovery questions and next step	34
13. Appendices	35

1. Executive summary

When an AI agent takes a consequential action, how can the enterprise prove that this exact action was authorized, remained safe at execution time, and produced only the intended effect?

Enterprises are moving from AI that *recommends* a decision to AI that *makes and executes* it. That shift is valuable, but it changes the governance question. A general, up-front approval establishes that an action of a certain kind is permitted; it does not, by itself, prove that **this exact action, on this exact target, was still safe and correct at the moment it ran** — or that the effect it produced matched what was authorized.

Three practical risks recur whenever an AI agent acts on the enterprise:

1. **The AI acts on the wrong target.** A valid-looking action is pointed at the wrong system, resource, or record.
2. **Conditions change between approval and execution.** A change freeze begins, a dependency degrades, or a cost limit shifts in the interval between approval and action.
3. **The recorded action and the actual operational effect do not match.** What the record says was authorized is not what actually happened in the environment.

What ServiceNow already provides

ServiceNow governs the enterprise AI estate and an increasingly capable agentic runtime. AI Control Tower discovers, governs, secures and measures AI; Action Fabric exposes a governed system of action to external agents; Agent Deviation Detection flags when an agent strays from its role; AI Risk & Compliance maps controls across frameworks; Change Management and the CMDB hold approvals, windows and business context; and — with NVIDIA OpenShell — central policy is enforced at runtime on file, command and network access. AI Control Tower also measures AI adoption, business impact, realized value and ROI. This briefing treats those capabilities as real strengths and builds on them.

What Ugence proposes to contribute

Ugence proposes a narrow, complementary layer: an **independent, verifiable authority-and-evidence artifact per business action**. For a consequential AI action it establishes, per request: whether a properly delegated authority approved it; that the action is bound to the exact target and payload; that required evidence is trusted and current; that the live environment is safe at execution time; and — afterward — that the observed effect matched the authorization. Every ServiceNow integration described here is **PROPOSED**; no ServiceNow connector ships today.

Why the relationship is complementary

ServiceNow remains the system of record, the system of action and the platform that enforces execution. Ugence contributes an independent authority-decision and evidence layer that **composes with** those capabilities. Ugence does not replace ServiceNow, does not sit above it, and does not claim unique cross-platform runtime enforcement — OpenShell already enforces at the kernel level across form factors. Where a Ugence differentiation is not established in ServiceNow documentation, this briefing states it as a *discovery question* for ServiceNow architects, never as a claim that ServiceNow lacks a capability.

Recommended bounded pilot

The recommended first pilot is **UC-5, autonomous change execution** — a controlled scaling workflow the customer already runs, where exact-target authorization, an independent live-safety clearance, and execution-to-effect reconciliation land cleanly on a small, reversible footprint. A 30–60-day bounded technical pilot is proposed in Section 11.

Standing disclaimer

All ServiceNow integrations described in this briefing are proposed. No ServiceNow connector currently ships. The scenarios are illustrative enterprise workflows, not claimed Ugence customer deployments.

2. What ServiceNow already provides

Positioning the proposal accurately begins with acknowledging ServiceNow's strengths. The capabilities below are current ServiceNow functionality, cited from ServiceNow-owned sources (see the source hierarchy in the appendices).

ServiceNow capability	What it does
AI Control Tower	Discovers, governs, secures, observes and measures AI across the enterprise, including real-time enforcement.
Action Fabric	A governed system of action over MCP/A2A; every action is identity-verified, permission-scoped and auditable.
Agent Deviation Detection	Flags at runtime when an agent strays from its authorized role.
AI Risk & Compliance	Multi-framework control mapping (EU AI Act, NIST AI RMF and others) and enforcement.
Change Management & CMDB	Approvals, change / blackout windows, conflict detection and business context.
Enterprise approvals & workflow	Platform-wide approval and workflow controls.
NVIDIA OpenShell	Central AI Control Tower policy enforced at runtime on file, command and network access.
Adoption, impact & ROI	Measures AI adoption, business impact, realized value and ROI.

Table 1. ServiceNow capabilities acknowledged in this briefing.

Framing discipline

This briefing does not state or imply that ServiceNow lacks governance, cannot enforce, or only inventories AI; that Ugence replaces ServiceNow or sits above it; or that Ugence uniquely provides cross-platform runtime enforcement. The relationship is complementary throughout.

3. Proposed ServiceNow–Ugence responsibility boundary

The division of responsibility is consistent across every scenario in this briefing and is the core of the partnership story. ServiceNow is the system of record, workflow and action platform; Ugence contributes an independent, signed authority-and-evidence layer that composes with it.

ServiceNow remains responsible for	Ugence proposes to contribute
Enterprise system of record	Binding decision artifacts
Workflow and case management	Trusted-evidence linkage
CMDB and business context	Exact-action and target authorization
Platform approvals and governance	Time- and scope-bounded authority
Action Fabric	Independent operational clearance
ServiceNow execution pathways	Governed execution coordination
AI Control Tower monitoring and controls	Execution receipts
Business-value and ROI reporting	Observed-effect reconciliation
	Authority lifecycle and reassessment
	Cross-stage evidence lineage
	Developing governed-value attribution

Table 2. Proposed responsibility boundary. The integration is **PROPOSED**; no ServiceNow connector currently ships.

Neither side replaces the other. Ugence composes with ServiceNow’s workflow, system of action and platform enforcement. The integration is marked **PROPOSED** throughout.

4. Understanding the data journey

The single most important idea in this briefing is that **not every Ugence module modifies the original ServiceNow business record**. Most modules read or reference business context and emit a separate governance artifact alongside it. Four kinds of data stay distinct:

Kind of data	Examples	Owner
1 · Original business data	Change, incident, CI, employee, entitlement, control evidence, amount, target system, requested action	ServiceNow (system of record)
2 · Derived governance artifacts	Decision record; model / action / risk authorization; clearance verdict; execution receipt; assurance result; revocation / reassessment signal	Ugence (each a separate artifact)
3 · Observed operational data	Infrastructure state, service health, blackout window, dependency condition, account status, execution result, post-action effect	ServiceNow + operational telemetry
4 · ServiceNow system-of-record entries	The change / incident / request record and its statuses; receives references, statuses and summarized outcomes	ServiceNow (via the proposed integration)

Table 3. Four kinds of data, kept distinct throughout.

Each Ugence module reads or references the necessary business context and emits a separate governance artifact. It does not silently rewrite the originating ServiceNow record.

Layman workflow

The same journey, in plain language. Under each stage is one short example of the information added.





Figure 1. The end-to-end data journey in plain language (illustrative).

Technical canonical workflow

The canonical technical spine. Each layer speaks a distinct decision verb, and authority never leaks across layers. Only the modules a given scenario actually uses are shown in that scenario's diagram; the full set is introduced in Section 5.



Where a scenario needs them, Policy Workflow Compiler and RA-5 sit upstream of the decision; Risk Authority mints a signed, scoped, time-limited authorization; Model Authority authorizes the model; StoryGraph flags sequence risk; RA-6 revokes authority; RA-7 assesses in-flight trajectory. Only participating modules are shown per scenario.

Figure 2. The canonical technical workflow, with the fail-closed stop path (illustrative).

5. Ugence module responsibilities

Each module is described at the interface level — what it receives, what it decides, the separate artifact it emits, and what it explicitly does not do. This briefing describes the guarantee each module provides, not the proprietary mechanism that produces it: no source code, complete schemas, algorithms, policy-compilation internals, canonicalization procedures, key handling or cryptographic implementation details are disclosed.

Architecture invariants

Verified responsibilities held throughout: the context-envelope record (CER) belongs to Decision Authority; Agent Runtime owns canonical execution state and coordinates governed execution (it is never labelled “Agent Runtime (CER)”; Agent Runtime invokes domain executors / providers; Cloud Scaling Operations is a domain executor, not an authorization gate; RA-8 compares authorized intent, execution evidence and observed effect, and never retroactively authorizes an action.

How to read the maturity labels

Each module card carries a package-maturity label defined in Appendix A. These labels describe the engineering status of the Ugence package only. REFERENCE-GRADE means a reference implementation exists and is offline-verifiable but is not production-validated; IMPLEMENTED means code exists and is verified at interface level. Neither implies production deployment validation, and neither implies a shipping ServiceNow integration — every ServiceNow integration in this briefing is separately marked PROPOSED, and no connector currently ships. Package availability, production validation and ServiceNow integration are three distinct dimensions and are never conflated.

5.1 Policy, evidence and decision

Policy Workflow Compiler

REFERENCE-GRADE

Question it answers	How is approved policy turned into deterministic, checkable constraints?
Receives / references	An approved, structured policy pack (compile-time, offline).
Checks / decides	Compiles policy into deterministic constraints.
Emits	A digest-addressed governed-workflow artifact.
Does not do	Make a binding decision; authorize, clear or run anything.

RA-5 Trusted Evidence Admission

REFERENCE-GRADE

Question it answers	Is each required control satisfied by trusted, re-checked evidence?
Receives / references	Control-evidence references.
Checks / decides	Whether evidence is trusted and current — a caller-asserted “pass” is inert.
Emits	An evidence-derived, re-checked control result.
Does not do	Add a second authority signature.

Decision Authority REFERENCE-GRADE	
Question it answers	Did a properly delegated authority — not the AI — approve this class of decision?
Receives / references	Proposed action, delegated authority, constraints; owns the context-envelope record (CER).
Checks / decides	Whether a binding decision exists within permitted scope, by a non-AI principal.
Emits	A separate, immutable decision record; owns execution and reconciliation records.
Does not do	Execute; inspect live conditions; let the AI self-authorize.

Risk Authority REFERENCE-GRADE	
Question it answers	How is an approved risk decision turned into a verifiable, bounded permission?
Receives / references	An allow-family risk decision.
Checks / decides	Whether to mint authority; scope never exceeds the decision.
Emits	A signed, scoped, time-limited authorization (tamper-evident).
Does not do	Execute the action.

5.2 Model, exact action, sequence and clearance

Model Authority REFERENCE-GRADE	
Question it answers	Which model may handle this specific request, right now?
Receives / references	Request context and approved model policy.
Checks / decides	Per-request eligibility.
Emits	ALLOW / DENY / HOLD / ESCALATE, with governed fallback and expiry.
Does not do	Execute the request; replace platform provider approval.

ActionGate REFERENCE-GRADE	
Question it answers	Is this the exact authorized action on the exact target, and is it still valid?
Receives / references	Decision record, target identity, proposed action.
Checks / decides	Whether the action presented matches the authorized action; uncertainty maps to INDETERMINATE.
Emits	AUTHORIZED / DENIED / INDETERMINATE, bound to the exact payload/target.
Does not do	Judge operational safety; own execution.

StoryGraph		REFERENCE-GRADE
Question it answers	Do individually-harmless steps add up to a harmful capability?	
Receives / references	A multi-step plan across steps.	
Checks / decides	Whether benign steps assemble a harmful capability.	
Emits	OBSERVE / ESCALATE (advisory).	
Does not do	Authorize or execute.	

Action Clearance (ACP)		REFERENCE-GRADE
Question it answers	Is it operationally safe to execute right now?	
Receives / references	An already-authorized action plus current live conditions.	
Checks / decides	Blackout, conflict, unhealthy dependency, operational hold.	
Emits	CLEAR / HOLD / BLOCK / ESCALATE (subtractive).	
Does not do	Create or broaden authority; dispatch execution.	

5.3 Coordination, actuation, lifecycle and assurance

Agent Runtime		REFERENCE-GRADE
Question it answers	How is the governed execution coordinated safely?	
Receives / references	A governed request; invokes providers / domain executors.	
Checks / decides	Execution lifecycle (retry / timeout / recovery).	
Emits	An execution receipt; owns canonical execution state.	
Does not do	Create authority; author policy; authorize; mint clearance.	

Cloud Scaling Operations	Core IMPLEMENTED · validation PILOT PENDING · additional capabilities UNDER DEVELOPMENT · integration PROPOSED	
Question it answers	How is a bounded infrastructure change carried out under strict controls?	
Receives / references	A scaling instruction plus an externally minted execution authorization.	
Checks / decides	Readiness; bounded execution (dry-run by default).	
Emits	An execution outcome + audit, returned to Agent Runtime.	
Does not do	Mint its own authority; act as a preceding authorization gate.	

RA-6 Authority Lifecycle**REFERENCE-GRADE**

Question it answers	How is authority kept current when conditions change?
Receives / references	A reassessment signal when conditions change.
Checks / decides	Whether to revoke / supersede / expire authority.
Emits	A lifecycle mutation; revocation bites at the next pre-effect recheck.
Does not do	Execute or authorize.

RA-7 Trajectory Assurance**REFERENCE-GRADE**

Question it answers	Is an in-flight execution drifting from expected behavior?
Receives / references	The in-flight execution (neutral observation).
Checks / decides	Whether the trajectory is drifting.
Emits	NORMAL / ESCALATED / UNKNOWN, a neutral reassessment signal.
Does not do	Mint authority.

RA-8 Execution Assurance**REFERENCE-GRADE**

Question it answers	Did reality match what was authorized?
Receives / references	Authorized action, execution receipt, observed post-action state.
Checks / decides	Whether only the intended target changed, within scope.
Emits	matched / mismatch / partial / unknown.
Does not do	Retroactively legitimize an unauthorized action.

6. Detailed lead scenario: UC-5 autonomous change execution

Illustrative scenario

This is an illustrative enterprise scenario showing how the proposed integration could work. It is not a claimed customer deployment.

Scenario at a glance:

- **ServiceNow change record:** CHG0048217
- **Business service:** online checkout
- **Target:** a production Kubernetes cluster
- **Proposed action:** scale the checkout service from 12 to 18 instances
- **Reason:** sustained demand and latency pressure
- **Constraints:** an approved cost ceiling, a permitted change window, healthy dependencies and rollback capability

6.1 The problem in plain language

Checkout latency is rising under sustained demand, and an AI operations agent recommends adding capacity. Scaling may prevent an outage — but the online-checkout service is revenue-critical, and so is any service that shares the cluster. Scaling the wrong cluster or service, or acting on stale information, could disrupt another service or exceed the approved cost. A valid approval may no longer be safe if a freeze has begun, a dependency has become unhealthy, or projected cost now exceeds the ceiling. And after the change, the record must reflect what actually happened — that only the intended instances changed, and nothing else. The enterprise wants the speed of autonomy together with proof that the exact approved change ran, at a safe moment, with a verifiable, reconciled result.

6.2 Data entering the workflow

Data category	Example	Authoritative source
Business record	Change request CHG0048217 and its approval state	ServiceNow Change Management
Target context	Cluster, checkout service and requested capacity (12 → 18)	CMDB, ITOM and cloud platform
Live conditions	Freeze window, dependency health and service health	ServiceNow and operational telemetry
Governance context	Delegated scaling authority, cost / risk limits and expiry	Approved enterprise policy
Intended outcome	Lower latency within cost and risk constraints	Business objective

Table 4. Business data entering the UC-5 workflow (business-readable fields, not internal schemas).

6.3 Layman workflow




Figure 3. UC-5 in plain language: the business journey (illustrative).

6.4 Technical module workflow

The workflow is presented in two phases for portrait reading: first authorization and clearance, then execution and assurance. The verified execution relationship is that Agent Runtime **invokes** Cloud Scaling Operations (a domain executor, not a preceding gate); the executor returns its result to Agent Runtime, which supplies the execution receipt to RA-8.

Phase 1 — Authorization and clearance





Figure 4. UC-5 Phase 1: binding decision, exact-action authorization and independent clearance, with the fail-closed stop path (illustrative).

Phase 2 — Execution and assurance



Agent Runtime invokes Cloud Scaling Operations (a domain executor, not a preceding gate); the executor returns its result to Agent Runtime, which supplies the execution receipt to RA-8. RA-8 compares authorized intent, the execution receipt and the observed effect; a mismatch / partial / uncertain result is escalated and still returns an auditable outcome to the change record. RA-8 never retroactively authorizes an action.

Figure 5. UC-5 Phase 2: Agent Runtime invokes Cloud Scaling Operations, receives its result, and supplies the execution receipt to RA-8 for effect reconciliation (illustrative).

6.5 Module-by-module walkthrough

Module	Receives	Question it answers	Adds	Sends next	If information is missing
Decision Authority	Proposed change, delegated authority, constraints	Is there a binding decision within delegated scope, by a non-AI principal?	A binding decision record	To ActionGate	No delegated authority → no binding decision; stop
ActionGate	Decision record, cluster identity, proposed action	Is this the exact authorized action, still valid?	An exact-target action authorization	To Action Clearance	Mismatch or uncertainty → INDETERMINATE; stop
Action Clearance	The authorization + current live conditions	Is it operationally safe right now?	A CLEAR / HOLD / BLOCK / ESCALATE verdict	To Agent Runtime (on CLEAR)	Unsafe or unknown live state → HOLD / BLOCK

Table 5. UC-5 authorization half — what each module receives, checks, adds and passes forward.

Module	Receives	Question it answers	Adds	Sends next	If information is missing
Agent Runtime	The governed request (on CLEAR)	How is execution coordinated safely?	An execution receipt; owns canonical execution state	Invokes Cloud Scaling Operations; carries the receipt to RA-8	No governance wired → fails closed; blocks
Cloud Scaling Operations	A scaling instruction + external execution authorization	Is it ready, and bounded to the authorization?	An execution outcome + audit	Returns the result to Agent Runtime	No execution authorization → dry-run only; no live change
RA-8 Execution Assurance	Authorized action, receipt, observed cluster state	Did only the authorized instances change, within scope?	matched / mismatch / partial / unknown	Status + receipt references to ServiceNow	Cannot verify → uncertain / manual review; never auto-authorize

Table 6. UC-5 execution half — coordination, actuation and assurance.

6.6 Human-control boundary

- **May be autonomous:** evaluating the approval, binding the exact action, checking live safety, and performing the bounded scaling inside the approved window and cost ceiling.
- **Forces HOLD / ESCALATE:** an active freeze, a conflicting change, an unhealthy dependency, projected cost over the ceiling, or an expired / superseded authorization.
- **Remains human-binding:** the delegation of scaling authority itself, and any change outside the approved class or scope.
- **Limits:** cost ceiling, capacity-delta bounds, blast radius (one named cluster / service) and reversibility (rollback requirement).

- **Fail-closed:** if a mandatory approval, evidence item, live-safety check or execution confirmation is missing, the workflow does not treat uncertainty as permission.

6.7 Business outcome and pilot measurements

Possible pilot measures include: exact-target authorization rate; authorization mismatches blocked; unsafe live conditions detected; execution-to-effect match rate; change success rate; rollback rate; latency improvement; service availability; time saved; and infrastructure cost.

Cloud Scaling maturity

Cloud Scaling maturity — four separate dimensions (do not conflate): Core Cloud Scaling Controller: IMPLEMENTED. Production validation: PILOT PENDING. Additional agentic-AI capabilities: UNDER ACTIVE DEVELOPMENT. ServiceNow integration: PROPOSED — no connector currently ships. The core controller is implemented and awaiting pilot validation; it is not downgraded because additional capabilities or pilot validation remain pending.

7. Additional detailed enterprise scenarios

7.1 UC-11 — Vulnerability remediation and emergency patching

A critical vulnerability is found on production servers, and an AI agent proposes to deploy an emergency patch. Speed matters — but patching the wrong set of servers, patching a business-critical system at the wrong moment, or continuing after the risk picture worsens mid-rollout can cause the very outage the patch was meant to prevent.


Illustrative: VUL0007731 → remediation task → change record; a critical remote-code-execution vulnerability on a web tier; 40 configuration items in “payments-web”; an emergency patch; a staged rollout with the ability to stop mid-flight. Not a customer deployment.

ServiceNow Vulnerability Response already creates remediation tasks, links them to change for approval and tracking, and can initiate emergency response workflows. The proposed Ugence value is exact-CI binding, per-stage clearance, mid-rollout revocation and effect reconciliation.

Data category	Example	Source
Business record	Remediation task and change record, approval state	ServiceNow Vulnerability Response + Change
Target context	40 CIs in payments-web, the patch	CMDB / Security Operations
Live conditions	Business-criticality, maintenance window, risk posture	ServiceNow + telemetry
Governance context	Delegated emergency-patch authority, expiry, revocability	Approved enterprise policy

Table 7. Business data entering the UC-11 workflow.

Layman workflow. Critical vulnerability found → AI proposes an emergency patch → approval confirmed for this exact server set and patch → permission bound to those servers → per-stage “is it safe right now?” → patch this stage (or stop remaining stages if unsafe or the risk changed) → confirm only authorized servers changed → auditable result returned to the remediation record.



Revocation is bounded-latency: it stops the next stage at the pre-effect recheck, not one already in progress. Illustrative example.

Figure 6. UC-11 technical workflow, with RA-6 mid-rollout revocation and the per-stage stop path (illustrative).

Human-control boundary. Per-stage clearance and patching are autonomous inside the authorized CI set and window; a peak-trading or criticality hold, a closed window, expired authority or a revocation signal forces HOLD / ESCALATE; the emergency-patch authority delegation and any CI outside the authorized set remain human-binding. Because RA-6 is the sole authority-lifecycle writer, a worsening risk posture can revoke authority mid-rollout; revocation is bounded-latency, stopping the *next* stage at the pre-effect recheck, not one already in progress.

Metrics and evidence. Mean-time-to-remediation; unauthorized-target attempts blocked; percentage of stages correctly stopped on revocation; effect-match rate. Evidence returned to ServiceNow: per-stage authorization, clearance, receipts, revocation record and effect-match.


7.2 UC-6 — Access provisioning with segregation of duties

An employee asks a self-service assistant for elevated access — for example, the ability to both create and approve payments. An AI agent could fulfil it in seconds, but granting an entitlement that breaks segregation of duties, or that the requester's current risk posture should block, can create fraud exposure that is hard to unwind.

Illustrative: RITM0102934; a finance analyst who already holds "create payment" requests "approve payment" in the ERP; a segregation-of-duties rule prohibits one person holding both. Not a customer deployment.

ServiceNow provides Service Catalog request fulfilment, entitlement management, and — with the Veza-based identity capabilities — least-privilege and AI Agent access controls. The proposed Ugence value is exact-entitlement authorization and an independent segregation-of-duties clearance.

Layman workflow. Employee requests elevated access → approval confirmed for this exact entitlement → permission bound to this user and entitlement → "any conflict? safe now?" → access granted, or blocked for human review on a segregation-of-duties conflict → confirm only that access was granted → auditable result returned to the request.



The decisive control is the segregation-of-duties conflict at clearance time: the action can be perfectly authorized and still be blocked. Illustrative example.

Figure 7. UC-6 technical workflow; the decisive control is the segregation-of-duties conflict at clearance time (illustrative).


Human-control boundary. Fulfilling entitlements with no segregation-of-duties conflict and a clear risk posture may be autonomous; any conflict, elevated risk posture, disabled account or expired authority forces HOLD / ESCALATE / BLOCK; the access-grant authority delegation and any entitlement outside scope remain human-binding. A perfectly valid authorization is still blocked when Action Clearance detects a conflict — authorization and clearance are independent.

Metrics and evidence. Access-policy violations prevented; percentage of grants with matched effect; time saved per request; unauthorized-entitlement attempts blocked. Evidence: decision, exact-entitlement authorization, SoD clearance and effect-match (or a clean “no change” on block).

7.3 UC-3 — High-risk AI action enforcement

Illustrative (privileged production access): an AI use case classified high-risk is about to grant privileged production access automatically. The policy is approved and controls are on file — but at the moment of action, is every required control currently satisfied by trusted, re-checked evidence, or is the system relying on a stale, self-asserted “pass”? Not a customer deployment.

Acting on out-of-date compliance evidence is exactly what audits punish. AI Control Tower and AI Risk & Compliance remain the governance systems of record — classification, multi-framework control mapping and enforcement. Ugence proposes an independently verifiable, action-level authority artifact linked to **current** evidence, so a stale or self-asserted control cannot wave an action through.



A caller-asserted “pass” is inert: only trusted, re-checked evidence satisfies a control. The authorization is signed, scoped and time-limited. AI Control Tower and AI Risk & Compliance remain the governance systems of record. Illustrative example.

Figure 8. UC-3 technical workflow: trusted-evidence admission, a signed and scoped authority, exact-action enforcement and effect reconciliation (illustrative).

Positioning (honest). Whether ServiceNow already emits an equivalent independently verifiable, evidence-fresh per-action artifact is a discovery hypothesis to confirm with the customer’s architects —

not an assumed gap. The policy approval and the authority delegation remain human-binding, and stale or self-asserted evidence is never treated as a satisfied control.


Metrics and evidence. Percentage of high-risk actions with fresh-evidence backing; stale-evidence blocks; effect-match rate. Evidence: a re-checked control result, a signed authorization, an exact-action authorization and an effect-match — an audit-ready chain.

7.4 UC-4 — External agents through Action Fabric, A2A and MCP

Enterprises now let external AI agents (for example Claude or Copilot) take real actions on enterprise systems. ServiceNow Action Fabric opens its system of action to these agents and routes every action through AI Control Tower; with NVIDIA OpenShell, policy is enforced at runtime on file, command and network access. The remaining question is narrow: for each business action an external agent takes, is there an independent, verifiable record that this exact action was authorized — and does anything notice when individually-allowed steps add up to something that should not be allowed?

Ownership

Ownership (unchanged): Action Fabric owns the governed ServiceNow action pathway; NVIDIA OpenShell owns its runtime / sandbox enforcement. Ugence proposes authority, exact-action binding, evidence lineage, clearance where applicable, sequence-risk and assurance artifacts. Ugence does not own or replace Action Fabric or OpenShell.



Dispatch and kernel-level enforcement are ServiceNow (Action Fabric) and NVIDIA (OpenShell). Ugence adds authority, exact-payload binding, sequence-risk and assurance only — composition, not replacement. Illustrative example.

Figure 9. UC-4 technical workflow: dispatch and kernel-level enforcement are ServiceNow (Action Fabric) and NVIDIA (OpenShell); Ugence adds authority, exact-payload binding, sequence-risk and assurance only (illustrative).

This is the strongest overlap zone and a partnership-native one — Anthropic is ServiceNow's Action Fabric design partner. Ugence claims neither a governance gap nor unique cross-runtime enforcement; the differentiation is confined to the properties of the independent authority-and-evidence artifact (Section 9). Metrics: unauthorized-payload attempts blocked; sequence escalations caught; percentage of actions with an independent authorization record.

8. Full use-case portfolio

The collection below comprises **enterprise workflow scenarios grounded in actual ServiceNow products, with proposed Ugence governance extensions**. They are not proven Ugence customer deployments. Every ServiceNow integration is **PROPOSED**.

8.1 Pilot-ready operational scenarios

UC-5 · Autonomous change execution **PROPOSED INTEGRATION**

Business problem	Scale / restart / config changes made autonomously
ServiceNow anchor	Change Management + ITOM
Ugence proposes	Exact-CI authority, independent clearance, effect reconciliation
Autonomy boundary	Autonomous within window and cost; freeze / conflict force stop
Discovery question	Which change and CMDB events are authoritative?

UC-6 · Access provisioning with SoD **PROPOSED INTEGRATION**

Business problem	Elevated access fulfilled by an assistant
ServiceNow anchor	Service Catalog / Employee Center (+ Veza)
Ugence proposes	Exact-entitlement authority + independent SoD clearance
Autonomy boundary	Autonomous when no SoD conflict; conflict → block
Discovery question	Which entitlement class suits a bounded pilot?

UC-11 · Vulnerability remediation **PROPOSED INTEGRATION**

Business problem	Emergency patch to production CIs
ServiceNow anchor	Security Operations — Vulnerability Response
Ugence proposes	Exact-CI binding, per-stage clearance, mid-rollout revocation
Autonomy boundary	Per-stage autonomy; revocation stops the next stage
Discovery question	Which critical-vulnerability workflow for a pilot?

8.2 Practical scenarios requiring bounded autonomy

UC-7 · Refunds and credits **PROPOSED INTEGRATION**

Business problem	AI issues a refund or credit
ServiceNow anchor	Customer Service Management (Now Assist)
Ugence proposes	Amount + account binding; reconcile executed = authorized
Autonomy boundary	Within a delegated dollar threshold
Discovery question	Which refund band is safe to pilot?

UC-8 · Procurement and PO execution**PROPOSED INTEGRATION**

Business problem	AI places a purchase order
ServiceNow anchor	Sourcing & Procurement Operations
Ugence proposes	Compiled policy, supplier-evidence check, exact PO binding
Autonomy boundary	Threshold-bound; above → human authority
Discovery question	Which spend threshold and category?

UC-10 · Agentic hiring (human-binding)**PROPOSED INTEGRATION**

Business problem	AI assists screening and scheduling
ServiceNow anchor	HR Service Delivery / Recruitment
Ugence proposes	Keep hire / reject human; immutable decision record
Autonomy boundary	Assistive only; decision stays human-binding
Discovery question	Which step may be automated vs human-bound?

8.3 Strategic governance scenarios

UC-2 · Per-request model authorization**PROPOSED INTEGRATION**

Business problem	Regulated data about to reach a model
ServiceNow anchor	AI Control Tower model / provider governance
Ugence proposes	Per-request ALLOW / DENY / HOLD / ESCALATE + fallback / expiry
Autonomy boundary	Per-request binding, not a static allowlist
Discovery question	Which data classes trigger per-request control?

UC-3 · High-risk action enforcement**PROPOSED INTEGRATION**

Business problem	High-risk AI action on stale evidence
ServiceNow anchor	AI Control Tower — AI Risk & Compliance
Ugence proposes	Signed, scoped, evidence-fresh action authority
Autonomy boundary	Only while trusted evidence is current
Discovery question	Which high-risk use case for a pilot?

UC-4 · External-agent governance**PROPOSED INTEGRATION**

Business problem	External agents act on the system of action
ServiceNow anchor	Action Fabric + AI Agent Fabric + OpenShell
Ugence proposes	Independent exact-payload authority, sequence-risk, assurance
Autonomy boundary	Composes with platform enforcement
Discovery question	Which external-agent workflow to pilot?

UC-12 · Data-boundary governance**PROPOSED INTEGRATION**

Business problem	Agents pull enterprise data into context
ServiceNow anchor	Workflow Data Fabric + AICT privacy
Ugence proposes	Minimum-necessary context; token accounting; fail-closed
Autonomy boundary	Governs exactly what crosses the model boundary
Discovery question	Which data domains need boundary control?

8.4 Emerging or future scenarios

UC-1 · Autonomous security containment**ANNOUNCED / FUTURE — Tier 2 SOC AI Specialist expected****December 2026**

Business problem	Autonomous host isolation / account disable / IP block
ServiceNow anchor	Security Incident Response — Tier 2 SOC AI Specialist
Ugence proposes	Exact-target authority, live-safety clearance, effect reconciliation
Autonomy boundary	Forward-looking; not shipped
Discovery question	When the specialist ships, where does authority bind?

UC-9 · Governed multi-agent workforce**PROPOSED INTEGRATION**

Business problem	A team of agents runs an end-to-end process
ServiceNow anchor	AI Agent Orchestrator / Autonomous Workforce
Ugence proposes	Least-privilege team plan granting nothing; per-action authorization + sequence-risk
Autonomy boundary	Composition grants no authority
Discovery question	Which multi-agent process to bound first?

9. Honest overlap and precise differentiation

Several capabilities genuinely overlap with current ServiceNow functionality. They are named here as overlaps, discussed as granularity or independent-verification extensions, never as governance gaps.

Zone	ServiceNow capability	Ugence proposed edge (narrow)
ActionGate ↔ Action Fabric / AICT	Governed, identity-verified, auditable action execution	Per-payload / target authorization re-checked at commit (discovery hypothesis)
RA-7 ↔ Agent Deviation Detection	Flags when an agent strays from its role	Independently verifiable trajectory signal into an authority lifecycle
RA-6 ↔ real-time shutdown	Can stop a misbehaving agent	Authority-lifecycle revoke / epoch; enforcement stays read-only
Model Authority ↔ Skill-Kit approval	Restricts which models run	Per-request binding decision with governed fallback and expiry
AICT + NVIDIA OpenShell	Central policy enforced at runtime across file / command / network	Ugence does not claim unique cross-runtime enforcement — the differentiation is the artifact

Table 8. Confirmed overlap zones and the narrow proposed Ugence edge.

The narrow, defensible differentiation is confined to the properties of an independent, action-level authority-and-evidence artifact: independently verifiable (by a party other than the executor); bound to the exact business action and target; linked to trusted evidence; scoped and time-limited; re-checked before consequential execution; separated from live operational clearance; connected to execution receipts; reconciled against observed effect; and usable in a cross-stage evidence lineage.

Evidence discipline

Absence from public documentation is never treated as proof that ServiceNow lacks a capability. Where a Ugence edge is unconfirmed, it is a discovery question for ServiceNow architects.

10. Enterprise Governed Value

Status: under development

Enterprise Governed Value is a developing, cross-cutting capability — not an authorization gate.

ServiceNow AI Control Tower already measures AI adoption, business impact, realized value and ROI. This proposal does not duplicate or replace that. The proposed Ugence contribution is **evidence-backed outcome attribution** connecting approved objectives, governed-execution evidence, attributable cost, observed outcomes, and preserved risk, compliance, quality and service constraints — reusing the governance evidence the pipelines already emit, adding attribution rather than a new gate.

Element	UC-5 example
Approved objective	Reduce checkout latency
Baseline	Current latency, capacity and cost
Governed action	Scale 12 → 18 instances
Execution evidence	Authorization · clearance · receipt
Observed outcome	Latency improved; service healthy
Attributable cost	Additional infrastructure expense
Preserved constraints	No change-window, risk or availability violation
Governed-value question	Enough attributable value without unacceptable trade-offs?

Table 9. Governed-value attribution illustrated with the UC-5 checkout-scaling example.

11. Proposed bounded pilot

The recommended first pilot is UC-5, autonomous change execution, as a bounded 30–60-day technical pilot. Autonomy is earned step by step; live controlled execution is not enabled until dry-run, exception-path and receipt checks pass. No broad production autonomy is promised at pilot start.

1. Joint technical discovery
2. Data and integration mapping
3. Historical replay
4. Dry-run or shadow evaluation
5. HOLD / BLOCK / ESCALATE testing
6. Limited controlled execution
7. Execution-receipt verification
8. Observed-effect reconciliation
9. Pilot evaluation
10. Joint next-phase decision

Possible pilot boundary: one controlled scaling workflow; a limited set of noncritical or carefully selected CIs / clusters; restricted action types; explicit authority and cost limits; defined rollback requirements; and no broad production autonomy at pilot start. Success is measured on safety and fidelity first — exact-target binding, mismatches blocked, effect-match — then on operational benefit.

12. Discovery questions and next step

Constructive questions to work through together — the answers shape the pilot:

- Which ServiceNow records are authoritative for each scenario?
- What events expose the approval, target and execution context?
- Where could an independent action-level authorization be evaluated?
- Where could a live operational clearance verdict be consumed?
- How should execution receipts and effect results be associated with the originating record?
- Which controls already provide equivalent granularity?
- What should remain human-binding?
- What data may leave the ServiceNow boundary?
- What evidence would be required for production acceptance?

The proposed next step is a joint technical discovery session to select and define one bounded pilot.

This is a request for technical fit, not a broad partnership commitment. ServiceNow remains the system of record, the system of action and the platform that enforces execution; Ugence contributes an independent, signed authority-and-evidence artifact that composes with it.

13. Appendices

Appendix A — Maturity and availability legend

Label	Meaning
IMPLEMENTED	Package code exists and is verified at interface level. Does not imply production-deployment validation or a shipping ServiceNow integration.
REFERENCE-GRADE	A reference implementation exists and is offline-verifiable, but is not production-validated. Does not imply a shipping ServiceNow integration.
PILOT PENDING	Implemented, but awaiting production / pilot validation.
ADDITIONAL CAPABILITIES UNDER DEVELOPMENT	Further capabilities beyond the implemented core are actively being built.
DESIGN-ONLY	Architecture / design intent; not yet a package.
PROPOSED INTEGRATION	A ServiceNow integration adapter is design intent only; no ServiceNow connector currently ships. (Also written “integration PROPOSED”.)
ANNOUNCED / FUTURE	A named future capability publicly announced but not yet shipped (e.g., the Tier 2 SOC AI Specialist expected December 2026).

Table 10. Maturity and availability labels used in this briefing.

Appendix B — Module-to-scenario matrix

Module	UC-5	UC-11	UC-6	UC-3	UC-4
Policy Workflow Compiler				•	
RA-5 Trusted Evidence Admission				•	
Decision Authority	•	•	•	•	
Risk Authority				•	
Model Authority					•
ActionGate	•	•	•	•	•
StoryGraph					•
Action Clearance (ACP)	•	•	•		•
Cloud Scaling Operations	•				
Agent Runtime	•	•	•	•	•
RA-6 Authority Lifecycle		•			
RA-7 Trajectory Assurance					•
RA-8 Execution Assurance	•	•	•	•	

Table 11. Which modules participate in each detailed scenario (• = participates).

UC-4 ends in RA-7 (in-flight trajectory assurance) rather than RA-8, because execution and enforcement are performed by ServiceNow Action Fabric and NVIDIA OpenShell. UC-11 uniquely exercises RA-6 for mid-rollout revocation. UC-5 uniquely exercises Cloud Scaling Operations as the domain executor.

Appendix C — Source hierarchy and citations

Product behavior is grounded first in ServiceNow product documentation, then release notes, Newsroom and Community; NVIDIA documentation for OpenShell; third-party press only for corroboration. Ugence module behavior is grounded in the source-of-record documents (catalog v1.2 and the use-case walkthrough companion) at the guarantee level; no internal repository detail is disclosed.

- **ServiceNow Docs:** existing product behavior and workflow mechanics (docs.servicenow.com).
- **ServiceNow release notes:** versions and availability.
- **ServiceNow Newsroom:** announcements, partnerships and future availability (e.g., the December 2026 Tier 2 SOC AI Specialist).
- **ServiceNow Community:** supporting and explanatory material.
- **NVIDIA documentation:** OpenShell technical behavior (enforcement via seccomp, Landlock LSM and network namespaces — not eBPF).
- **Third-party press:** corroboration of announcement facts only.

Direct page fetches to some ServiceNow properties were egress-restricted during preparation; product names, availability framing and protocol support were corroborated across multiple ServiceNow-owned sources and should be spot-checked on the live documentation before customer use. No internal repository paths are cited in this external document.

Appendix D — Standing qualifications and disclaimers

- All ServiceNow integrations described are PROPOSED; no ServiceNow connector currently ships.
- All scenarios are illustrative enterprise workflows, not claimed Ugence customer deployments.
- Maturity dimensions are not conflated: a multi-module workflow carries per-module maturity.
- Unverified differentiation is a discovery hypothesis, never a claim that ServiceNow lacks a capability.
- This briefing describes what each module guarantees, not the proprietary mechanism that produces it.