

Ugence Developer Agentic Studio Explainer

One entry per screen, with the screen as captured, and now for each screen exactly what to enter, what to press and what to expect on the live deployment.

The three public surfaces, and what each one serves today (2026-09-12). Governance Studio <https://studio-web-production-65e8.up.railway.app> · Governance Console <https://console-web-production-4e77.up.railway.app> · Authority Plane <https://authority-plane-production.up.railway.app>. The screen captures in this document were taken on 2026-09-07 from a local composition in which every seam was wired: a fixture provider, a policy registry, a review service, a worker with a seeded directory (tenant `tenant-a`). The Railway deployment is thinner, by the documented start commands (`docs/deployment/RAILWAY_REFERENCE_DEPLOYMENT.md`): the studio API starts bare, so every optional v2 seam is absent and screens 14, 15, 16, 19, 20, 21, 22, 23, 24 and 25 answer with a typed *Not available in this deployment* notice `[I]`; screens 1 to 13, 17 (Validate), 18 and 33 work in full. The console runs its three scenarios in shadow `[V]`. The authority plane reads an empty directory in tenant `tenant-demo`: Grants and Holders answer 200 with zero rows, Committee and Grant events answer a typed 404 `[V]`. The URLs could not be opened from the environment that produced this revision, so every live-state claim is inferred from the deployment reference and marked `[I]`; open `/studio/status` and the console's Modules tab to confirm before a demo.

Contents

How to use this document on the live deployment

The one thing to know first

Module coverage

A · Governance Studio, Eligibility Explorer (13 screens, contract v1)

B · Governance Studio, Governed Agent Studio (12 screens, contract v2)

C · Console (3 tabs, the packaged console API)

D · Authority Plane (4 screens, the governed runtime worker's reads)

E · Governance Studio, Bring Your Workflow (1 screen, contract v1)

What is not a screen, and why

Status: reference, 2026-09-07; revised 2026-09-12 with an *Enter, press, expect* block per screen and the live-deployment state. Screen 33 revised 2026-09-13 for the Bring Your Workflow phase 3A draft controls (BW-3A). Written after the owner's ruling that AP-3 controls (`ADR_UGENCE_AUTHORITY_PLANE_SCOPING.md` §18) and the Bring Your Workflow ruling (§22). Every entry is taken from the screen's own source: its stated purpose, its stated disclaimer, the operations it calls, and the ruling that shaped it. Nothing here describes a screen as doing more than its code does.

Four front ends, thirty-three screens. Each entry answers the same five questions: **what it shows**, **who answers it**, **what the operator can do**, **what it never does**, and **the ruling that shaped it**.

How to use this document on the live deployment

Each screen entry now ends with an **Enter, press, expect** block: the exact values to type or select, the button to press, and the output those values produce, with the source that fixes them. Three rules apply throughout.

- **Studio explorer routes are nested under the scenario.** Screens 3 to 13 live at

`/scenarios/<scenario_id>/<section>` (for example `/scenarios/procurement/eligibility`), not at the flat paths the headings carry. Only `/scenarios`, `/scenarios/:id` and `/bring-your-workflow` are top-level `[V]`

(frontend/src/app/App.tsx:46-61).

- **Four scenario ids exist and no others:** `procurement`, `customer_support`, `cybersecurity_success`, `cybersecurity_no_feasible_team` `[V]` (backend/.../scenarios/catalog.py:36-41). Any other id is a 404.

- **A typed gap is a correct answer, not a fault.** On the live studio most Governed Agent Studio screens report *Not available in this deployment* with the seam they lack. Read the notice aloud; it is the deployment saying what it does not have rather than showing a green tick over nothing.

The demo order that works is the console first, then the studio, then the plane (docs/deployment/CLIENT_DEMO_PLAN.md).

The one thing to know first

None of these is an administration screen. Ruling MA-1 refused a per-module admin screen; ruling AP-1 to AP-5 admitted an authority plane instead, and at this step that plane reads and does not write. Across all four front ends there are exactly two kinds of screen: **reads** (display what a package, service or deployment returned) and **typed intakes** (record what a person asserted, conferring nothing). The acts that would change what an agent may do (issue, activate, revoke, grant, authorize, clear, execute) are performed on no screen anywhere, by ruling SD-2. The worker implements loading and revoking a role grant behind its identity gate, but under AP-3, which the owner confirmed as controlling (§18), neither is served until the identity adapter is validated end to end against a real enterprise issuer; the in-process issuer evidence is implementation and conformance evidence only.


Module coverage

MODULE	WHERE IT REACHES A SCREEN
Context Minimization, Truth Assurance, ActionGate, Autonomous Control Plane	Console: probes on Modules, stages of the Governed Loop, entries in Audit
Agent Runtime	Studio: Simulate (the in-studio path over fixtures), and the worker's shadow run started from Simulate and watched from Review
Agent Workforce Composer's adapter over an operator's own document	Studio: Bring Your Workflow (validate, adapt, compare; since BW-3A also keep as an unapproved DRAFT, list, read back, supersede)
Model Selection, Hybrid LLM, LLM Steering, Autonomous Runtime	No screen. By ruling MS-1 their registry rows do not reach the studio either
The authority directory and approval workflow (not among the nine modules)	Authority Plane: all four screens; Studio: Review Queue and Run Detail

A · Governance Studio, Eligibility Explorer (13 screens, contract v1)

Every screen here reads the frozen `governance_studio.api.v1` contract, twenty approved operations over the Agent Workforce Composer, against pinned synthetic scenarios. Nothing on these screens is editable except the What-If controls, which act on a temporary copy. The shared banner on every page says it: eligibility is not selection, assignment, authorization or execution.

1 Scenario catalog (/scenarios)



Screen 1, as captured 2026-09-07

- **Shows:** the pinned synthetic scenarios, each with its title and maturity.
- **Answered by:** `list_scenarios`.
- **Operator can:** open a scenario.
- **Never:** creates, edits or imports a scenario.
- **Ruling:** P3B/P3C screen audit; the frozen v1 contract.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>https://studio-web-production-65e8.up.railway.app/scenarios</code> .	Four cards: Procurement Sourcing Workforce (badged <i>Recommended demo</i>), Customer Support Triage & Response, Cybersecurity Incident Response (Feasible), Cybersecurity Incident Response (No Feasible Team). Each shows <i>Synthetic data</i> , <code>Workflow contract = workflow_ir.v1</code> and its expected state.
2	Press Open scenario on the procurement card.	The overview at <code>/scenarios/procurement</code> .

Values you can type

- There is nothing to type on this screen. The expected states are `COMPLETE` for three scenarios and `NO_FEASIBLE_TEAM` for the fourth.

Source: `frontend/src/features/scenarios/ScenarioCatalog.tsx:8-76`;
`backend/.../scenarios/catalog.py:36-66`.

2 Scenario overview (/scenarios/:id)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

⚠ Eligibility ≠ Selected ≠ Assigned ≠ Authorized ≠ Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

Procurement Sourcing Workforce

Non-greedy agent team selection under provider concentration limits.
non-greedy team selection under provider concentration

Eligibility means an agent remains permitted for consideration under the pinned workflow, registry and enterprise policies. It does **not** mean the agent was selected, assigned, authorized or executed.

9Workflow nodes

8Workflow edges

3AI-agent roles

6Non-agent steps

5Registered agents

4Eligible pairs

11Ineligible pairs

0Indeterminate pairs

IDENTITY & VERSIONS

Scenario ID	procurement
Domain	procurement
Workflow identity	gs_procurement
Workflow contract	workflow_ir.v1
API contract	governance_studio.api.v1
AWC version	0.2.1
Compiler version	consumed via AWC adapter
Fixture version	0.2.0

DETERMINISTIC VERIFICATION

✔ Observed eligibility fingerprint matches the frozen expected value

Expected fingerprint	c19735272e08...61e900
Observed fingerprint	c19735272e08...61e900
Workflow adaptation	91c1f7d2533b...bce25a

Screen 2, as captured 2026-09-07

- **Shows:** metadata, versions, verification, digests, maturity and presentation counts.
- **Answered by:** `get_scenario`, `get_scenario_workflow`, `get_scenario_registry`, `get_scenario_eligibility`, `get_version`.
- **Operator can:** read, and navigate to the other screens.
- **Never:** shows a ranking or team-selection metric.
- **Ruling:** screen audit §12.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement</code> .	Identity block: Scenario ID <code>procurement</code> , Workflow identity <code>gs_procurement</code> , contract <code>workflow_ir.v1</code> , fixture version <code>0.2.0</code> . Tiles: 9 nodes, 8 edges, 3 AI-agent roles, 6 non-agent steps, 5 agents, 4 eligible pairs, 11 ineligible, 0 indeterminate.

Ugence Developer Agentic Studio Explainer, with inputs · 2026-09-13

8 / 78


STEP	DO THIS	EXPECT
2	Read the <i>Deterministic verification</i> line.	Green check: <i>Observed eligibility fingerprint matches the frozen expected value</i> (sha256: e179463...a2833).
3	Open <code>/scenarios/cybersecurity_no_feasible_team</code> .	7 nodes, 6 edges, 2 roles, 4 agents, 2 eligible, 6 ineligible; headline names the infeasibility.

Values you can type

- customer_support: 6 nodes, 5 edges, 3 roles, 6 agents, 7 eligible, 11 ineligible.
- cybersecurity_success: 9 nodes, 8 edges, 4 roles, 6 agents, 5 eligible, 19 ineligible.

Source: `expected_outputs/<id>/adaptation.json`, `eligibility.json`;
`demo_data/<id>/scenario_manifest.json`; `ScenarioOverview.tsx:65-106`.

3 Workflow (/workflow)



Screen 3, as captured 2026-09-07

- **Shows:** the compiled workflow as a graph and an accessible list, synchronised with a node details panel; complete node and edge accounting; every disposition as the API returned it.
- **Answered by:** `get_scenario_workflow`.
- **Operator can:** select nodes, zoom, fit.
- **Never:** edits the workflow or computes a disposition in the browser.
- **Ruling:** screen audit §13 to §15.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/workflow</code> .	Header <code>9 nodes · 8 edges · contract workflow_ir.v1</code> ; the first node is pre-selected.
2	Click node <code>proc_supplier_evidence</code> in the graph or the list.	Details panel: disposition <code>AI_AGENT_ELIGIBLE</code> , link View role requirements .
3	Click <code>proc_binding_approval</code> .	Disposition <code>HUMAN_AUTHORITY_REQUIRED</code> ; no role link.

Values you can type

- procurement dispositions: `proc_request_validation` DETERMINISTIC_SERVICE_PREFERRED;
`proc_supplier_evidence`, `proc_supplier_risk`, `proc_recommendation` AI_AGENT_ELIGIBLE;
`proc_binding_approval` HUMAN_AUTHORITY_REQUIRED; `proc_purchase_auth`, `proc_commit_clearance`
EXISTING_GOVERNANCE_CAPABILITY_OWNS_STEP; `proc_audit`
DETERMINISTIC_SERVICE_PREFERRED; `proc_terminal` NO_AI_AGENT_REQUIRED.

Source: `expected_outputs/procurement/adaptation.json`; `WorkflowScreen.tsx:29-33`;
`NodeDetails.tsx:75-80`.

4 Role (/roles/:roleId)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

🚩 Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

< All scenarios

procurement recommendation

AI-agent role

recommendation_draft

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

IDENTITY

Role ID

role::proc_recommendation

Source node

proc_recommendation

Workflow identity

gs_procurement

Role fingerprint

098123ab3dc1_bf172f

FUNCTIONAL REQUIREMENTS

Required capabilities

evidence_extraction, procurement_recommendation

Optional capabilities

None

Domain

None

Supported tools

None

Input contracts

supplier_risk_report

Output contracts

recommendation_draft

ENTERPRISE CONSTRAINTS

Provider rules

None

Enterprise policy

Residency

None

Deployment

None

Data classification

Not supplied

Security floor

0

Audit requirement

None

Evidence classes

MEASURED

Quality floor

Not supplied

Latency ceiling

Not supplied

Cost ceiling

Not supplied

AUTHORITY & PERMISSION BOUNDARIES

These are permission requirements used as eligibility inputs. P3C does not show composition-time proposal bundles, feasibility scoring or the granting of permissions (P3D).

Required permissions

read_context

Prohibited permissions

None

Authority ceiling

0

Human review

false

PROVENANCE

Compiler contract

awc.v1

AWC-derived

Source package digest

179d78b83963_2811e8

Policy references

None

Screen 4, as captured 2026-09-07

- **Shows:** one role's requirements, each field badged with its source: compiler, enterprise policy or composer.
- **Answered by:** `get_scenario_workflow`.
- **Operator can:** read.
- **Never:** infers a field.
- **Ruling:** screen audit §16.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open /scenarios/procurement/roles/role::proc_supplier_evidence (type the double colon literally).	Heading <i>supplier evidence collection</i> , badge <i>AI-agent role</i> ; required capabilities <code>evidence_extraction</code> , <code>supplier_evidence_collection</code> badged <i>Compiler</i> ; source node <code>proc_supplier_evidence</code> .
2	Type a wrong id, e.g. /scenarios/procurement/roles/role::nope.	Empty state: <i>Role not found</i> · <i>No AI-agent role role::nope in this scenario</i> .

Values you can type

- procurement: `role::proc_supplier_evidence`, `role::proc_supplier_risk`, `role::proc_recommendation`.
- customer_support: `role::sup_triage`, `role::sup_retrieval`, `role::sup_draft`.
- cybersecurity_success: `role::sec_evidence_collection`, `role::sec_threat_analysis`, `role::sec_incident_correlation`, `role::sec_recommendation`.
- cybersecurity_no_feasible_team: `role::sec_threat_analysis`, `role::sec_incident_correlation`.

Source: `expected_outputs/<id>/adaptation.json` → `role_requirements`; RoleScreen.tsx:22-116.

5 Agent registry (/registry)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

🛑 Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed ✖

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

< All scenarios

5 synthetic agents · snapshot 7068ec47aeeb...d4f545

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

Agent registry

agent_general_analyst

v1.0.0 · openal

Synthetic

ACTIVE

Domain	procurement, support, security		
Residency	US		
Deployment	cloud		
Security class	3		
Tools	web_search		
Requested permissions	read_context		
Max authority	1		
Audit capabilities	trace, replay		
Validity	0 → Not supplied		
Profile fingerprint	—		
DECLARED EVIDENCE 0			
None			
MEASURED EVIDENCE 2			
Capability	Value	Benchmark	It
evidence_extraction	0.95 score	bench_evidence_extraction	s
supplier_evidence_collection	0.9 score	bench_supplier_evidence_collection	s
OBSERVED EVIDENCE 0			
None			

agent_india_procurement

v1.0.0 · anthropic

Synthetic

ACTIVE

Domain	procurement, support, security		
Residency	IN		
Deployment	cloud_in		
Security class	3		
Tools	web_search		
Requested permissions	read_context		
Max authority	1		
Audit capabilities	trace, replay		
Validity	0 → Not supplied		
Profile fingerprint	—		
DECLARED EVIDENCE 0			
None			
MEASURED EVIDENCE 3			
Capability	Value	Benchmark	It
evidence_extraction	0.95 score	bench_evidence_extraction	s
procurement_risk_analysis	0.95 score	bench_procurement_risk_analysis	s
supplier_evidence_collection	0.95 score	bench_supplier_evidence_collection	s
OBSERVED EVIDENCE 0			
None			

agent_procurement_recommendation

Synthetic

ACTIVE

agent_procurement_risk

Synthetic

ACTIVE

Screen 5, as captured 2026-09-07

- **Shows:** the pinned synthetic agent registry; evidence grouped as DECLARED, MEASURED or OBSERVED, always labelled synthetic; expired evidence stays visible and flagged.
- **Answered by:** `get_scenario_registry`.
- **Operator can:** read.
- **Never:** hides expired evidence or registers an agent.
- **Ruling:** screen audit §17.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/registry</code> .	Header 5 synthetic agents · snapshot sha256:7068ec4...d545. Five agent cards, evidence grouped DECLARED / MEASURED / OBSERVED, each badged Valid or Expired.

STEP	DO THIS	EXPECT
2	Find <code>agent_india_procurement@1.0.0</code> .	Provider <code>anthropic</code> , residency <code>IN</code> , deployment <code>cloud_in</code> ; the one non-US agent, which the eligibility screen eliminates.

Values you can type

- procurement agents: `agent_general_analyst@1.0.0` (openai, US), `agent_india_procurement@1.0.0` (anthropic, IN), `agent_procurement_recommendation@1.2.0`, `agent_procurement_risk@2.1.0`, `agent_supplier_evidence@1.4.0` (anthropic, US).
- cybersecurity_success includes `agent_low_clearance@1.0.0` at security level 2, eliminated on every role.

Source: `demo_data/<id>/agent_registry_snapshot.json`; `RegistryScreen.tsx:13-103`.

6 Eligibility matrix (/eligibility)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

Eligibility ✕ Selected ✕ Assigned ✕ Authorized ✕ Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

Eligibility matrix

5 role-agent pairs · 1 eligible · 4 ineligible · 0 indeterminate

Eligibility means an agent remains permitted for consideration under the pinned workflow, registry and enterprise policies. It does **not** mean the agent was selected, assigned, authorized or executed.

procurement recommendation · supplier evidence collection · procurement risk analysis

Provider

Residency

Agent status

Elimination reason

Sort by

All

All

All

All

Agent identity

Reset

Display order is not a selection decision. No agent is ranked or preferred.

Agent	State	P / F / U	agent_status	agent_version	audit_capabil...	authority_cei...	capability_ev...	capabilit
agent_general_analyst openai · v1.0.0	✕ Ineligible	19/1/0	✓	✓	✓	✓	✓	·
agent_india_procurement anthropic · v1.0.0	✕ Ineligible	17/3/0	✓	✓	✓	✓	✓	·
agent_procurement_recommendation anthropic · v1.2.0	✓ Eligible	21/0/0	✓	✓	✓	✓	✓	✓
agent_procurement_risk anthropic · v2.1.0	✕ Ineligible	19/1/0	✓	✓	✓	✓	✓	·
agent_supplier_evidence anthropic · v1.4.0	✕ Ineligible	19/1/0	✓	✓	✓	✓	✓	·

Result fingerprint (first visible row): 885e4fa520a5_75657d

Synthetic demonstration data · deterministic planning only · no agent execution · no permission granting · no business-action authorization

Screen 6, as captured 2026-09-07

- **Shows:** agents by rows, API-provided condition names by columns, every state from the API; an explanation drawer per cell.
- **Answered by:** `get_scenario_eligibility`, `explain_eligibility`.
- **Operator can:** filter, reset filters, open an explanation.
- **Never:** shows a score, rank, recommendation or preference.
- **Ruling:** screen audit §18 to §20; the primary P3C screen.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/eligibility</code> .	Role <i>supplier evidence collection</i> selected; header 5 <i>role-agent pairs</i> · 2 <i>eligible</i> · 3 <i>ineligible</i> · 0 <i>indeterminate</i> ; 21 condition columns.
2	Set Provider to <code>anthropic</code> .	Four rows remain.

STEP	DO THIS	EXPECT
3	Set Elimination reason to <code>RESIDENCY_MISMATCH</code> .	One row: <code>agent_india_procurement@1.0.0</code> , <code>INELIGIBLE</code> , failed <code>residency</code> and <code>deployment</code> .
4	Press Explain on that row.	Drawer: reason <i>Residency does not satisfy the requirement</i> , evidence refs <code>ev::agent_india_procurement::...:MEASURED</code> , policy refs <code>gs_procurement_enterprise_policy</code> , <code>gs_eligibility_policy</code> , both fingerprints.
5	Press Reset , then choose role <i>procurement recommendation</i> .	1 eligible (<code>agent_procurement_recommendation@1.2.0</code>), 4 ineligible.

Values you can type

- Filters are selects only; over-filtering shows *No agents match the filters*.
- `customer_support`: `agent_threat_analysis@1.0.0` is ineligible on all three roles (`MISSING_REQUIRED_CAPABILITY`).
- `cybersecurity_success`: `agent_low_clearance@1.0.0` fails `security_classification` on every role.

Source: `expected_outputs/procurement/eligibility.json`; `EligibilityScreen.tsx:75-205`;
`ExplanationDrawer.tsx:84-126`.

7 Ranking (/ranking)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

🚩 Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed

Synthetic demonstration data · Deterministic planning only · No live agent execution · No permission granting · No business-action authorization · Not pilot validated · Not production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-If

Ranking

role ranking fingerprint 268861cf99bd...ff8282

Ranking, composition, permission proposals and fallback plans are deterministic planning outputs over synthetic data. They do not assign agents, grant permissions, provision runtime access, authorize actions or execute workflows.

role::proc_recommendation

role::proc_supplier_evidence

role::proc_supplier_risk

Presentation order

Canonical rank (API)

Presentation order does not change the canonical API rank. 1 eligible · 4 excluded

Rank	Agent	Provider	Total score	Tie group	Score breakdown
1	agent_procurement_recommendation v1.2.0	anthropic	7589	1	Show breakdown

Synthetic demonstration data · deterministic planning only · no agent execution · no permission granting · no business-action authorization

Screen 7, as captured 2026-09-07

- **Shows:** the canonical API rank order, score decomposition and tie-break as the API gave them.
- **Answered by:** `get_scenario_ranking` , `explain_ranking` .
- **Operator can:** switch to canonical order.
- **Never:** computes a score in the browser.
- **Ruling:** P3D §12 to §13.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/ranking</code> ; role <code>role::proc_supplier_evidence</code> .	Rank 1 <code>agent_supplier_evidence@1.4.0</code> total 7819; rank 2 <code>agent_general_analyst@1.0.0</code> total 7604; caption 2 <i>eligible</i> · 3 <i>excluded</i> .

STEP	DO THIS	EXPECT
2	Press Show breakdown on rank 1.	Eight criteria in order: evidence_strength, evidence_freshness, measured_quality, observed_reliability, latency_headroom, cost_efficiency, security_headroom, audit_strength; tie-break vector begins 7819, 5000, 9090 .
3	Change Presentation order to <i>identity</i> , then press Reset to canonical rank .	Rows reorder, then return to API order.

Values you can type

- customer_support sup_triage: agent_support_triage 7859 over agent_multilingual_support 7579.
- cybersecurity_no_feasible_team: one candidate per role, 7889 and 7869.

Source: expected_outputs/<id>/ranking.json; RankingScreen.tsx:55-152 .

8 Composition (/composition)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.ap1.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

Composition

Complete

Ranking, composition, permission proposals and fallback plans are deterministic planning outputs over synthetic data. They do not assign agents, grant permissions, provision runtime access, authorize actions or execute workflows.

Role	Selected primary	Score	Top-ranked	Non-greedy?	Assignment fingerprint
role::proc_recommendation	★ Selected primary agent_procurement_recommendation v1.2.0	7589	agent_procurement_recommendation	No	618d3bfc48f ac9514
role::proc_supplier_evidence	★ Selected primary agent_general_analyst v1.0.0	7604	agent_supplier_evidence	Yes	c6adf6562e2 01912c
role::proc_supplier_risk	★ Selected primary agent_procurement_risk v2.1.0	7639	agent_procurement_risk	No	0d8123e3415 3378fa

ROLE RANKING VS. TEAM COMPOSITION

Candidate ranking evaluates role-level suitability. Team composition applies additional workflow-wide constraints and objectives, so the selected team may differ from independently choosing each role's top-ranked candidate. A lower-ranked selected candidate is not a mistake.

role::proc_supplier_evidence: top-ranked agent_supplier_evidence · selected agent_general_analyst — see team constraints/objectives below.

TEAM-LEVEL FACTS

Unfilled rolesNone

Total team score26832

HARD CONSTRAINTS

max_roles_per_agent

satisfied (1/2)

provider_concentration

satisfied (2/3/67%)

failure_domain_concentration

satisfied (2/3/67%)

TEAM OBJECTIVES

aggregate_ranking_quality

contribution 22832

provider_diversity

contribution 2000

failure_domain_diversity

contribution 2000

Screen 8, as captured 2026-09-07

- **Shows:** plan state, assignments, team-level facts, the non-greedy explanation, distinct selection states, and an honest NO_FEASIBLE_TEAM view.
- **Answered by:** `get_scenario_plan`, `explain_plan`.
- **Operator can:** read.
- **Never:** reruns the composition search.
- **Ruling:** P3D §14 to §17.

Enter, press, expect

On the live deployment: works in full.

Ugence Developer Agentic Studio Explainer, with inputs · 2026-09-13

20 / 78

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/composition</code> .	Plan COMPLETE, id <code>plan::a403d4bf91a614b5</code> , team score 26832; assignments <code>role::proc_recommendation → agent_procurement_recommendation@1.2.0</code> , <code>role::proc_supplier_evidence → agent_general_analyst@1.0.0</code> , <code>role::proc_supplier_risk → agent_procurement_risk@2.1.0</code> .
2	Read the <i>Non-greedy</i> column.	Yes only on <code>role::proc_supplier_evidence</code> : the top-ranked agent (7819) was not chosen, to keep provider concentration at 2 of 3 under the 67% limit. Optimality EXACT_OPTIMUM, search space 2, feasible teams 1.
3	Open <code>/scenarios/cybersecurity_no_feasible_team/composition</code> .	The NO_FEASIBLE_TEAM card: no assignments, score 0, feasible teams 0, termination <i>exhaustive bounded search complete</i> . Both roles are fillable only by anthropic agents against a 60% concentration limit and minimum provider diversity 2.

Values you can type

- This screen has no controls.

Source: `expected_outputs/<id>/agent_team_plan.json`;
`demo_data/cybersecurity_no_feasible_team/composition_policy.json`; `CompositionScreen.tsx:118-237`.

9 Permission proposals (/permissions)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

EligibilitySelectedAssignedAuthorizedExecuted

Synthetic demonstration dataDeterministic planning onlyNo live agent executionNo permission grantingNo business-action authorizationNot pilot validatedNot production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

Permission proposals

Ranking, composition, permission proposals and fallback plans are deterministic planning outputs over synthetic data. They do not assign agents, grant permissions, provision runtime access, authorize actions or execute workflows.

Permission proposals are planning artifacts. They do not grant, provision or activate permissions, and they are not authorizations.

role::proc_recommendation

agent_procurement_recommendation · v1.2.0

Feasible

PROPOSED PERMISSIONS (LEAST-PRIVILEGE)

read_context

CATEGORIZED 1

read_context

Proposed

Authority scope1

NoticeThis is a planning-time permission-bound proposal. It does not grant, authorize, provision or execute any permission.

Proposal fingerprint4900f2f023cb..53ab73

role::proc_supplier_evidence

agent_general_analyst · v1.0.0

Feasible

PROPOSED PERMISSIONS (LEAST-PRIVILEGE)

read_context

CATEGORIZED 1

read_context

Proposed

Authority scope1

NoticeThis is a planning-time permission-bound proposal. It does not grant, authorize, provision or execute any permission.

Proposal fingerprintf4a463b6acbd..709b9a

role::proc_supplier_risk

agent_procurement_risk · v2.1.0

Feasible

PROPOSED PERMISSIONS (LEAST-PRIVILEGE)

read_context

CATEGORIZED 1

read_context

Proposed

Authority scope1

NoticeThis is a planning-time permission-bound proposal. It does not grant, authorize, provision or execute any permission.

Proposal fingerprintf4a463b6acbd..709b9a

Screen 9, as captured 2026-09-07

- **Shows:** composition-time permission proposals, categorised, with feasibility.
- **Answered by:** `get_scenario_plan`.
- **Operator can:** read.
- **Never:** describes a proposal as granted, provisioned, active or authorized. A grant is an authority act, and no screen performs one (SD-2).
- **Ruling:** P3D §18 to §19.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/permissions</code> .	Three cards, one per assigned role, each <i>Feasible</i> , proposed permission <code>read_context</code> categorised <code>PROPOSED</code> , notice <i>This is a planning-time permission-bound proposal...</i>

STEP	DO THIS	EXPECT
2	Open /scenarios/cybersecurity_no_feasible_team/permissions.	Empty state: <i>No permission proposals.</i>

Values you can type

- No controls.

Source: `PermissionScreen.tsx:19-36`; `expected_outputs/<id>/agent_team_plan.json`.

10 Fallbacks (/fallbacks)

Ugence Governance StudioEligibility Explorer

Governed Agent Studiogovernance_studio.api.v1

🛑 Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed ·
Synthetic demonstration data · Deterministic planning only · No live agent execution · No permission granting · No business-action authorization · Not pilot validated · Not production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-if

Fallbacks

Ranking, composition, permission proposals and fallback plans are deterministic planning outputs over synthetic data. They do not assign agents, grant permissions, provision runtime access, authorize actions or execute workflows.

3

Roles with primaries

1

Roles with ≥1 fallback

2

Roles with no fallback

1

Fallback candidates

role::proc_recommendation

primary agent_procurement_recommendation · v1.2.0

✖ No fallback available

This role has no independent fallback under the pinned registry and policies.

Fallback fingerprint394aad3a1e6b_c9a787

role::proc_supplier_evidence

primary agent_general_analyst · v1.0.0

🔔 Limited fallback

ORDERED FALLBACK CANDIDATES 1

#1 agent_supplier_evidence v1.4.0rank #1, score 7819bp, diverse failure domain

Fallback fingerprint7e11a8479d0c_9ae449

role::proc_supplier_risk

primary agent_procurement_risk · v2.1.0

✖ No fallback available

This role has no independent fallback under the pinned registry and policies.

Fallback fingerprinta279eed7c492_dfb745

Synthetic demonstration data · deterministic planning only · no agent execution · no permission granting · no business-action authorization

Screen 10, as captured 2026-09-07

- **Shows:** every role's fallback coverage or gap; the summary counts API-returned states only.
- **Answered by:** `get_scenario_plan`.
- **Operator can:** read.
- **Never:** selects a fallback.
- **Ruling:** P3D §20 to §21.

Enter, press, expect

On the live deployment: works in full.


STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/fallbacks</code> .	Tiles 3 roles / 1 complete or partial / 2 without fallback / 1 candidate. <code>role::proc_supplier_evidence</code> PARTIAL with <code>agent_supplier_evidence v1.4.0</code> (rank #1, score 7819bp, diverse failure domain); the other two roles <code>NO_FALLBACK_AVAILABLE</code> .
2	Open <code>/scenarios/customer_support/fallbacks</code> .	<code>role::sup_retrieval</code> COMPLETE with two candidates; the other two PARTIAL.

Values you can type

- No controls.

Source: `FallbackScreen.tsx:26-42; expected_outputs/<id>/agent_team_plan.json.`

11 Replay (/replay)



Screen 11, as captured 2026-09-07

- **Shows:** deterministic reconstruction of the plan and whether it matched; a mismatch is a prominent integrity state; a deterministic export.
- **Answered by:** `replay_plan`, `export_scenario`.
- **Operator can:** run the replay, export the scenario.
- **Never:** reruns agent execution or suppresses a mismatch.
- **Ruling:** P3D §22 and §26.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/replay</code> .	Runs on load. Green check: <i>Plan replayed deterministically (fingerprints match)</i> ; expected and replayed fingerprint both <code>sha256:c197352...e900</code> .
2	Press Load export manifest , then Download export bundle .	<code>governance-studio-procurement-export.json</code> is saved locally.

Values you can type

- Plan fingerprints: customer_support sha256:16ca014...e8ae8 , cybersecurity_success sha256:9be9cf6...4be98 , cybersecurity_no_feasible_team sha256:e4d7ebd...f147b .
- A mismatch would render *REPLAY MISMATCH — integrity check failed*; it cannot occur on the frozen data.

Source: ReplayScreen.tsx:54-92; demo_data/<id>/scenario_manifest.json → expected_fingerprints.plan_fingerprint.

12 Plan comparison (/compare)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

EligibilitySelectedAssignedAuthorizedExecuted

Synthetic demonstration dataDeterministic planning onlyNo live agent executionNo permission grantingNo business-action authorizationNot pilot validatedNot production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-If

Plan comparison

Ranking, composition, permission proposals and fallback plans are deterministic planning outputs over synthetic data. They do not assign agents, grant permissions, provision runtime access, authorize actions or execute workflows.

Right-hand plan

Forbidden provider

Baseline with a provider forbidden

anthropic

Both plans are produced deterministically by the API; the diff is API-computed.

Request failed (403)

request failed with status 403

Synthetic demonstration data - deterministic planning only - no agent execution - no permission granting - no business-action authorization

Screen 12, as captured 2026-09-07

- **Shows:** the API's diff of two deterministically produced plans from the same scenario.
- **Answered by:** `compare_plans`.
- **Operator can:** choose the two plans.
- **Never:** diffs in the browser.
- **Ruling:** P3D §23.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>/scenarios/procurement/compare</code> .	Right-hand plan defaults to <i>Baseline with a provider forbidden</i> , provider <code>anthropic</code> .
2	Leave the defaults.	Same workflow Yes; Plan A <code>sha256:c197352...</code> , Plan B <code>sha256:6eebd94...11c6</code> ; 3 assignment changes, each ... -> <code>None</code> ; 3 permission and 3 fallback changes; policy digest change <code>enterprise</code> ; score delta <code>-26832</code> .

STEP	DO THIS	EXPECT
3	Switch to <i>Identical baseline (control)</i> .	<i>The two plans are identical — no assignment, constraint, permission, fallback or policy change.</i>

Values you can type

- Provider options: `anthropic`, `openai` (plus `google` on `cybersecurity_no_feasible_team`). There are no free plan ids; the left plan is always the scenario baseline.

Source: `CompareScreen.tsx:17-88`; `frontend/tests/fixtures/procurement.compare.json`.

13 Controlled what-if (/what-if)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

EligibilitySelectedAssignedAuthorizedExecuted

Synthetic demonstration dataDeterministic planning onlyNo live agent executionNo permission grantingNo business-action authorizationNot pilot validatedNot production certified

< All scenarios

Procurement Sourcing Workforce

Overview

Workflow

Registry

Eligibility

Ranking

Composition

Permissions

Fallbacks

Replay

Comparison

What-If

Controlled what-if

Ranking, composition, permission proposals and fallback plans are deterministic planning outputs over synthetic data. They do not assign agents, grant permissions, provision runtime access, authorize actions or execute workflows.

What-if analysis evaluates a temporary copied scenario. It does not modify the frozen scenario, enterprise policy, registry or any production system.

Perturbation (bounded)

provider

FORBID_PROVIDER

anthropic

Apply

Synthetic demonstration data · deterministic planning only · no agent execution · no permission granting · no business-action authorization

Screen 13, as captured 2026-09-07

- **Shows:** the modified plan and diff the API returned for one bounded perturbation.
- **Answered by:** scenario_what_if.
- **Operator can:** pick one of nine allowlisted operations with constrained controls and apply it.
- **Never:** mutates the frozen scenario; the API evaluates a temporary copy. No arbitrary JSON, policy, URL or code input.
- **Ruling:** P3D §24 to §25, C2.

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open /scenarios/procurement/what-if.	Perturbation defaults to FORBID_PROVIDER, provider anthropic.

STEP	DO THIS	EXPECT
2	Press Apply .	State change <code>COMPLETE</code> → <code>NO_FEASIBLE_TEAM</code> ; 3 assignment, 3 permission, 3 fallback changes; baseline and modified fingerprints differ; <i>What-if analysis evaluates a temporary copied scenario...</i>
3	Press Reset to baseline ; choose <code>REMOVE_CANDIDATE</code> , candidate <code>agent_supplier_evidence@1.4.0</code> ; Apply .	Plan stays <code>COMPLETE</code> ; a fallback change on <code>role::proc_supplier_evidence</code> (its only fallback is gone).
4	Choose <code>TIGHTEN_PROVIDER_CONCENTRATION</code> , type <code>25</code> ; Apply .	Plan becomes <code>NO_FEASIBLE_TEAM</code> .
5	Type <code>150</code> in the same field.	Inline error <i>limit_pct must be <= 100</i> ; Apply is disabled.

Values you can type

- The nine operations and a working procurement value each: `FORBID_PROVIDER anthropic`; `REQUIRE_RESIDENCY IN` (the seeded default; drives the plan infeasible); `TIGHTEN_COST_CEILING 1.0`; `TIGHTEN_LATENCY_CEILING 100.0`; `REVOKE_AGENT_VERSION agent_general_analyst@1.0.0`; `EXPIRE_EVIDENCE` (no parameter; logical time jumps to 3 000 000); `TIGHTEN_PERMISSION_POLICY invoke_tool`; `TIGHTEN_PROVIDER_CONCENTRATION 25`; `REMOVE_CANDIDATE agent_supplier_evidence@1.4.0`.
- Numeric fields start empty and are required; ceilings must be finite and ≥ 0 ; the concentration limit must be an integer 0 to 100.

Refusals to expect

- ceiling is required, must be a finite number, ceiling must be ≥ 0 , limit_pct must be an integer, limit_pct must be ≤ 100 .*

Source: `features/whatif/operations.ts:47-184`; `WhatIfScreen.tsx:81-149`;
`backend/.../services/orchestration.py:206-288`; `frontend/e2e/what-if-matrix.spec.ts:22-32`.

B · Governance Studio, Governed Agent Studio (12 screens, contract v2)

Every screen here reads the additive `governance_studio.api.v2` contract, twenty-six operations, amended seven times by sha256-chained record. The studio backend reaches each governance package only through the eleven-entry SD-1 allowlist. The banner on every page says no screen here issues, activates, revokes, grants, authorizes, clears or executes, and a test enforces it on every operation id, path and summary.

14 Registration (/studio/registration)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.ap1.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.ap1.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Registration

Record what an administrator asserts about one AI system: its exact binding, an owner reference, a classification label and a validity window.
Registering confers nothing. This screen never admits, approves, gates, promotes, attests, edits, revokes or deletes; a changed system is a new registration that supersedes the old.

Register a system

Typed fields only. The tenant is this deployment's, the registration id is derived by the registry, and the owner reference is recorded as PRESENTED_UNPROVEN: an opaque handle, not an authenticated identity. The classification label is recorded uninterpreted.

Binding id *

Subject id *

Context id *

Context digest (sha-256 hex) *

System id *

System version *

Configuration id *

Configuration digest (sha-256 hex) *

Deployment environment ref

Owner reference *

Classification label *

Issued at (ISO-8601, with timezone) *

Expires at (ISO-8601, with timezone)

Supersedes (registration id)

Notes

Register

Registrations in force

Screen 14, as captured 2026-09-07

- **Shows:** the systems registered in this deployment's one tenant, as of an instant.
- **Answered by:** `v2_registry_list`, `v2_registry_register`; the `ai-system-registry` package over a sqlite file under the runtime volume.
- **Operator can:** record what an administrator asserts about one AI system: its binding, an owner reference, a classification label and a validity window.
- **Never:** admits, approves, gates, promotes, attests, edits, revokes or deletes; a changed system is a new registration that supersedes the old. Registering confers nothing.
- **Ruling:** front-door seam 5, FD-9; the amendment v2-A1.

Enter, press, expect

On the live deployment: typed gap `system_registry`: no system registry is configured: this deployment holds no registration file, so nothing can be recorded or listed `[I]`. The values below succeed where a registry and the classification vocabulary are composed.

Ugence Developer Agentic Studio Explainer, with inputs · 2026-09-13

33 / 78

STEP	DO THIS	EXPECT
1	Fill the nine binding fields: Binding id <code>bind-1</code> ; Subject id <code>subject-1</code> ; Context id <code>ctx-1</code> ; Context digest 64 lowercase hex (for example <code>64 × a</code>); System id <code>hiring-screener</code> ; System version <code>1.0.0</code> ; Configuration id <code>cfg-1</code> ; Configuration digest 64 × <code>a</code> ; Deployment environment ref blank.	No error while typing; validation happens on Register.
2	Owner reference <code>directory://people/owner-1</code> ; Classification label <code>HIGH_RISK</code> ; Issued at <code>2026-09-01T00:00:00+00:00</code> ; Expires at <code>2027-09-01T00:00:00+00:00</code> ; Notes <code>first</code> .	Ready to submit.
3	Press Register .	<i>Registered reg_... · owner reference PRESENTED_UNPROVEN · confers nothing; registry_kind SqliteSystemRegistry, a derived 32-hex registration_id, record_digest; the record appears under Registrations in force with as_of_source request.</i>
4	Press Register again with the same values.	Refused <code>registration_duplicate</code> .

Values you can type

- Classification vocabulary `eu-ai-act-system-classification 1.0.0: PROHIBITED, HIGH_RISK, TRANSPARENCY_OBLIGATIONS, MINIMAL_RISK, UNCLASSIFIED`. The label is recorded uninterpreted; an unknown label is not refused.
- Digests must match `^[0-9a-f]{64}$`; instants must be ISO-8601 with a timezone; there is no tenant, registration id or registered-by field to type.

Refusals to expect

- `registration_refused` *AssessedSystemBinding.context_digest must be a lowercase 64-char sha-256 hex digest; validity.issued_at must be an ISO-8601 instant with a timezone; SystemRegistration.owner_ref must be a non-empty string; supersession_refused* when Supersedes names a record that is not this system's predecessor.

Source: `RegistrationScreen.tsx:20-147`; `services/studio_v2.py:853-1006`; `backend/tests/test_registry_seam.py:38-86`; vocabulary `docs/vocabularies/eu-ai-act-system-classification/1.0.0.json`.

15 Data use (/studio/data-use)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data Deterministic planning only No live agent execution No permission granting No business-action authorization Not pilot validated Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Data use

Record what a declarer asserts about the data one AI system uses: an opaque reference to it, what they called it, what they said it is for, and for how long.

A declaration confers nothing. This screen never inspects, classifies, redacts, minimizes, admits, authorizes, verifies, scores or enforces; it restricts no egress, and a changed declaration is a new one that supersedes the old.

Declare a data use

Typed fields only. The tenant is this deployment's, the declaration id is derived by the package, and the declarer is recorded as PRESENTED_UNPROVEN: an opaque handle, not an authenticated identity. The data reference is a locator and never the data. The classification, purpose and residency labels are recorded exactly as typed and interpreted nowhere.

Binding id *

Context id *

System id *

Configuration id *

Deployment environment ref

Classification label *

Recorded uninterpreted: what the declarer called it, never what that means.

Residency label

Recorded as metadata and evaluated nowhere in this deployment.

Expires at (ISO-8601, with timezone)

Supersedes (declaration id)

Subject id *

Context digest (sha-256 hex) *

System version *

Configuration digest (sha-256 hex) *

Data reference *

An opaque locator — a dataset id, a record handle. Never the data itself.

Purpose label *

Recorded uninterpreted.

Issued at (ISO-8601, with timezone) *

Declared by

An opaque handle, recorded as presented and unproven.

Notes

Screen 15, as captured 2026-09-07

- **Shows:** the data-use declarations of this tenant, as of an instant.
- **Answered by:** `v2_data_use_list`, `v2_data_use_declare`; the `data-use-admission` package.
- **Operator can:** record what a declarer asserts about the data one system uses: an opaque reference, its name, its purpose, for how long.
- **Never:** inspects, classifies, redacts, minimizes, admits, authorizes, verifies, scores or enforces; restricts no egress.
- **Ruling:** front-door seam 8, FD-12; v2-A4.

Enter, press, expect

On the live deployment: typed gap `data_use_declarations`: *no data-use declarations file is configured* `[I]`.

STEP	DO THIS	EXPECT
1	Fill the binding as on screen 14 (digests 64 × <code>b</code>).	—

STEP	DO THIS	EXPECT
2	Data reference <code>dataset://applicants/2026</code> ; Classification label <code>CONFIDENTIAL</code> ; Purpose label <code>shortlisting</code> ; Residency label <code>eu-west</code> ; Issued at <code>2026-09-01T00:00:00+00:00</code> ; Expires at <code>2027-09-01T00:00:00+00:00</code> ; Declared by <code>directory://people/declarer-1</code> .	Ready to submit.
3	Press Declare .	<code>declared true</code> , <code>store_kind</code> <code>SqliteDataUseDeclarations</code> , a derived <code>decl_...</code> id, <code>declared_by_status</code> <code>PRESENTED_UNPROVEN</code> , <code>confers</code> <code>nothing...</code> , <code>egress_restrictions</code> <code>not expressible...</code> ; the declaration is listed.

Values you can type

- Classification vocabulary `data-classification` 1.0.0: `PUBLIC` , `INTERNAL` , `CONFIDENTIAL` , `RESTRICTED` .
Purpose vocabulary `data-use-purpose` 1.0.0 is an open shape: any non-empty token is admitted.
- The data reference is an opaque locator, never the data; residency is recorded and evaluated nowhere.

Refusals to expect

- `declaration_refused` on a blank data reference or purpose, a label with control characters, a naive instant; `declaration_duplicate` ; `supersession_refused` .

Source: `DataUseScreen.tsx:24-233` ; `studio_v2.py:1014-1161` ; `backend/tests/test_data_use_seam.py:55-69` .

16 Vendor dependencies (/studio/vendor)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.ap1.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.ap1.v2

additive contract alongside the frozen v1 explorer

planning, preflight and observation only

no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Vendor dependencies

Record what a declarer asserts about one AI system's dependency on a vendor: an opaque reference to the vendor, what they called its risk posture, the policy they cite, and for how long.

A declaration confers nothing. This screen never resolves, verifies, scores, grades, ranks, approves, onboards or contacts; the risk posture is recorded, never assessed, and a changed declaration is a new one that supersedes the old.

Declare a vendor dependency

Typed fields only. The tenant is this deployment's, the declaration id is derived by the package, and the declarer is recorded as PRESENTED_UNPROVEN: an opaque handle, not an authenticated identity. The vendor reference is a locator and never a way to reach the vendor. Nothing here is a supplier record, a questionnaire or a due-diligence workflow.

Binding id *

Context id *

System id *

Configuration id *

Deployment environment ref

Risk posture label *

Issued at (ISO-8601, with timezone) *

Declared by

An opaque handle, recorded as presented and unproven.

Notes

Subject id *

Context digest (sha-256 hex) *

System version *

Configuration digest (sha-256 hex) *

Vendor reference *

An opaque locator in your own spelling. Never an address, endpoint or credential.

Policy reference *

Recorded and never resolved: a reference that names nothing looks the same here as one that names a real policy.

Expires at (ISO-8601, with timezone)

Supersedes (declaration id)

Screen 16, as captured 2026-09-07

- **Shows:** the vendor-dependency declarations of this tenant.
- **Answered by:** `v2_vendor_list`, `v2_vendor_declare`; the `vendor-dependency` package.
- **Operator can:** record an opaque vendor reference, a risk posture as declared, a cited policy, a window.
- **Never:** resolves, verifies, scores, grades, ranks, approves, onboards or contacts; the risk posture is recorded, never assessed.
- **Ruling:** front-door seam 9, FD-13; v2-A5.

Enter, press, expect

On the live deployment: typed gap `vendor_declarations`: no vendor declarations file is configured `[I]`.

STEP	DO THIS	EXPECT
1	Fill the binding as on screen 14 (digests 64 × e).	—

Ugence Developer Agentic Studio Explainer, with inputs · 2026-09-13

37 / 78

STEP	DO THIS	EXPECT
2	Vendor reference <code>vendor://acme-llm</code> ; Risk posture label <code>ASSESSED_WITH_FINDINGS</code> ; Policy reference <code>policy://vendor-standard/v3</code> (required); Issued at <code>2026-09-01T00:00:00+00:00</code> ; Expires at <code>2027-09-01T00:00:00+00:00</code> ; Declared by <code>directory://people/declarer-1</code> .	Ready to submit.
3	Press Declare .	<code>declared true</code> , a derived id, <code>risk_posture</code> <code>uninterpreted: ... ordered, compared, ranked and scored nowhere</code> , <code>confers nothing...</code> .

Values you can type

- Posture vocabulary `vendor-dependency-assessment-state 1.0.0: NOT_ASSESSED`, `ASSESSED_NO_FINDINGS`, `ASSESSED_WITH_FINDINGS`, `ASSESSMENT_LAPSED`, `ASSESSMENT_REFUSED`. The policy reference is text and is never resolved.

Refusals to expect

- `vendor_declaration_refused` on a blank vendor or policy reference; `vendor_declaration_duplicate`; `supersession_refused`.

Source: `VendorScreen.tsx:23-225`; `studio_v2.py:1207-1344`; `backend/tests/test_vendor_seam.py:51-65`.

17 Constitution (/studio/constitution)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data · Deterministic planning only · No live agent execution · No permission granting · No business-action authorization · Not pilot validated · Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 · additive contract alongside the frozen v1 explorer · planning, preflight and observation only · no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Constitution

Author a constitution, check its structure, and dry-run every pre-signing check.
This screen never issues and never activates a constitution — those are authority acts.

Constitution document

JSON document

```
{
  "constitution_id": "con_demo",
  "clauses": [],
  "version": "1"
}
```

Validate

Preflight issuance

Synthetic demonstration data · deterministic planning only · no agent execution · no permission granting · no business-action authorization

Screen 17, as captured 2026-09-07

- **Shows:** the structural validation of an authored constitution and a dry run of every pre-signing check.
- **Answered by:** `v2_constitution_validate`, `v2_constitution_preflight`; the constitution policy and activation packages over the activation root handed at composition.
- **Operator can:** author, validate, preflight.
- **Never:** issues or activates a constitution; those entry points are permanently outside the studio's allowlist. Preflight reports a typed gap where no signing key or trust root exists.
- **Ruling:** SD-2; front-door seam 1, FD-5.

Enter, press, expect

On the live deployment: **Validate** works in full. **Preflight issuance** answers the typed gap `constitution_preflight`: *no ActivationRoot is configured: this repository ships no signing key and no trust root* `[I]`.

STEP	DO THIS	EXPECT
1	Press Validate on the pre-filled sample <code>{"constitution_id":"con_demo","clauses": [],"version":"1"}</code> .	<code>validation_state INVALID</code> , diagnostic <code>invalid_constitution</code> . The sample is not a constitution artifact; the screen shows that honestly.

STEP	DO THIS	EXPECT
2	Replace the document with the minimal valid constitution listed below and press Validate .	<code>validation_state</code> <code>VALID</code> , no diagnostics, <code>constitution_id</code> <code>ugence.agent-constitution/tenant-1/baseline/v1</code> , a canonical digest.
3	Reorder <code>permitted_candidate_dispositions_bound</code> to put <code>RECOMMEND_MATCHED_FOR_APPROVAL</code> first and press Validate .	INVALID: lists must be in ascending codepoint order; the package refuses rather than reordering.
4	Press Preflight issuance .	Live: the <code>constitution_preflight</code> gap. With a root composed the button still returns <code>REFUSED</code> / <code>approval_reference_unstructured</code> , because the screen sends no approval reference; a report needs a direct POST (see values).

Values you can type

- Minimal valid document: `{"metadata":{"policy_id":"agent-constitution-baseline","version":"1.0.0","content_digest":"<64 lowercase hex>","scope":"TENANT","lifecycle_state":"APPROVED_ACTIVE","tenant_id":"tenant-1","supersedes_ref":"","supersedes_coordinate":null,"effective_from":"2026-01-01T00:00:00Z","effective_to":"2027-01-01T00:00:00Z"},"agent_constitution_ref":"ugence.agent-constitution/tenant-1/baseline/v1","governed_role_refs":["ugence.roles/tenant-1/proposer-reviewer/v1","ugence.roles/tenant-1/reconciler/v1"],"permitted_candidate_dispositions_bound":["ESCALATE_EXCEPTION","RECOMMEND_MATCHED_FOR_APPROVAL"],"permitted_review_actions_bound":["CREATE_EXCEPTION_REVIEW_BUNDLE"],"permitted_tool_scopes_bound":["scope.evidence-read","scope.report-write"],"constitution_vocabulary_version":"ugence.agent-constitution/clauses/v1"}`
- `scope` is `GLOBAL` (`tenant_id` must be empty) or `TENANT` (`tenant_id` required); `lifecycle_state` ∈ `DRAFT`, `APPROVED_ACTIVE`, `SUPERSEDED`, `WITHDRAWN` and only `APPROVED_ACTIVE` is issuable.
- Candidate dispositions vocabulary: `ESCALATE_EXCEPTION`, `RECOMMEND_MATCHED_FOR_APPROVAL`, `RECOMMEND_WITHHOLD`, `REQUEST_EVIDENCE`. Review actions: `CREATE_EXCEPTION_REVIEW_BUNDLE`, `ROUTE_APPROVAL_BUNDLE`. Every list unique and sorted; tool scopes may be empty.
- A preflight report (where a root exists) needs `POST /api/v2/constitution/preflight` with `approval_reference` shaped `<approving_authority_id> + | + <approval_ref> + | + <64 hex>`; rows artifact-recognition, reference-tenant, supersession, body-digest (passes only if the declared `content_digest` equals the canonical body digest), lifecycle, effectivity pass, and the approval row fails because the deployment composes a deny-all verifier.

Refusals to expect

- Not valid JSON*: ... on a parse error; `invalid_constitution` for any structural fault, including an unsorted or duplicate token, a non-hex digest, a naive datetime, or a `GLOBAL` scope with a tenant.

Source: `ConstitutionScreen.tsx`:18-113; `studio_v2.py`:162-230; `packages/integration/agent-constitution-policy/.../policy.py`:144-410; `backend/tests/test_constitution_real_root.py`:64-135.

18 Policy (/studio/policy)

Ugence Governance Studio Eligibility Explorer Governed Agent Studio governance_studio.ap1.v1

⚠ **Eligibility** ✖ **Selected** ✖ **Assigned** ✖ **Authorized** ✖ **Executed** ✖
Synthetic demonstration data Deterministic planning only No live agent execution No permission granting No business-action authorization Not pilot validated Not production certified

Registration Data use Vendor dependencies Constitution **Policy** Authority Simulate Publish Observe Review Status

governance_studio.ap1.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Policy

Author a policy pack on the canvas, preview the Workflow IR it compiles to, and compile a reviewed pack.
This screen never grants a permission and never executes a workflow; the compiler decides what is valid.

Canvas

Capability CONNECTOR_MAPPING Role AUTHORITY_REQUIREMENT Obligation AUDIT_REQUIREMENT Policy clause DECISION_RULE

9 governance objects across four kinds. The canvas owns capabilities, roles, obligations and policy clauses; every other part of the pack passes through unchanged.

Screen 18, as captured 2026-09-07

- **Shows:** a policy pack on the canvas, the Workflow IR it compiles to, and the result of compiling a reviewed pack.
- **Answered by:** `v2_policy_validate`, `v2_policy_synthesize`, `v2_policy_compile`; the policy-workflow compiler.
- **Operator can:** author on the canvas, validate, preview, compile. Compile requires a human approval record the compiler validates.
- **Never:** grants a permission or executes a workflow; the compiler decides what is valid. Today the screen works on a frozen fixture pack; a typed pack intake is a candidate for a later ruling.
- **Ruling:** GAS-R3, the canvas as the ratified authoring surface.

Enter, press, expect

On the live deployment: works in full; the compiler runs in-process.

STEP	DO THIS	EXPECT
1	Read the canvas.	9 governance objects in four kinds: 3 capabilities (supplier id, budget id, amount mappings), 2 roles (human purchase approver; budget authorization), 3 obligations (decision, action authorization, execution/reconciliation audits), 1 policy clause (purchase approval decision). Pack <code>pack.procurement.reference</code> , status APPROVED, version 1.
2	Press Validate .	A validation report with <code>policy_pack_id</code> <code>pack.procurement.reference</code> and no blocking diagnostics.
3	Press Preview Workflow IR .	<code>synthesized true</code> ; the nodes of the compiled workflow (the same nine nodes screen 3 shows for procurement).
4	Press Compile with approval .	<i>Compiled sha256:fb9fd4b9...8158a</i> ; the assurance manifest; the session now holds a release that screen 21 can send. The digest is frozen and identical on every call.

Values you can type

- There is nothing to type: the pack and the approval record are the frozen fixtures `fixtures/v2/policy_pack.json` and `approval_record.json`.
- The approval record the compile requires: id `approval.procurement.reference.fixture`, decision APPROVED, pack digest `sha256:28eedf60...98b14`, reviewer `fixture.reviewer` (`procurement_governance_reviewer`), approved 2026-08-03, `is_fixture true`.

Refusals to expect

- Omitting the approval (API only) is a 422; a compiler refusal renders *The compiler refused to synthesize this pack* with the exception name.

Source: `PolicyScreen.tsx:18-136`; `studio_v2.py:282-306`; `expected_outputs/v2_policy_compile.json`; `backend/tests/test_v2_routes.py:96-125`.

19 Authority (/studio/authority)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

⚠ Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Authority

Which policies are issued, resolvable, revoked or superseded, and the decisions a run rested on.

This screen is a reader. It never issues, revokes or supersedes a policy.

Issued policy records

Not available in this deployment: authority_registry

no PolicyRegistry is configured: the only reachable implementation is in-memory and holds one process's view, so an empty list would misrepresent an enterprise registry

Synthetic demonstration data - deterministic planning only - no agent execution - no permission granting - no business-action authorization

Screen 19, as captured 2026-09-07

- **Shows:** which policies are issued, resolvable, revoked or superseded, and the decisions a run rested on; which registry answered, with a warning when it is in-memory.
- **Answered by:** `v2_authority_list_policies`, `v2_authority_read_policy`, `v2_authority_read_decision`; the policy-authority registry handed at composition; the decision store is absent by ruling and reports its gap.
- **Operator can:** read.
- **Never:** issues, revokes or supersedes a policy.
- **Ruling:** front-door seam 2, FD-6.

Enter, press, expect

On the live deployment: typed gap `authority_registry`: *no PolicyRegistry is configured: the only reachable implementation is in-memory...* `[I]`.

STEP	DO THIS	EXPECT
1	Open the screen.	One read, no controls. Live: the gap notice above.

STEP	DO THIS	EXPECT
2	Where a sqlite policy registry is composed (compose profile).	<i>No issued records for the identities this deployment queries</i> plus <code>registry_kind SqlitePolicyRegistry</code> and <code>identities_queried</code> such as <code>agent_governance.agent_constitution agent-constitution-baseline TENANT tenant-1</code> . The registry ships empty; issuance is outside the studio's allowlist, so nothing ever appears here from a screen.

Values you can type

- No inputs. A policy identity, where configured, is `<policy_family>|<policy_id>|<scope>` with the tenant appended by the composition.

Source: `AuthorityScreen.tsx:12-56`; `studio_v2.py:381-460`; `deployment/governance-studio/.../app.py:93-101`.

20 Simulate (/studio/simulate)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.ap1.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.ap1.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Simulate

Two distinct paths: run a compiled workflow in this studio against fixtures, or ask the governed runtime worker to start its own shadow run. Each is labelled with its executor, its hook and its maturity. Nothing consequential is reachable from here. There is no live execution mode on either path, and no path is chosen for you.

Path A - In-process fixture run

Executor this studio's own Agent Runtime (seam 3)

Hook the runtime's default hook: every consequential task blocks (GOVERNANCE_NOT_CONFIGURED)

Maturity DEMONSTRATION_ONLY fixture provider; no external effect

Execution mode

DRY_RUN

Run simulation

Path B - Worker shadow run

Executor the governed runtime worker, a separate unit (seam 6); the studio executes nothing

Hook the worker's governed hook over the approval-bound source: a consequential task parks on ESCALATE in the review queue until a recorded human decision

Maturity REFERENCE_GRADE_SHADOW_ONLY; the worker's providers are FIXTURE_ONLY

The worker holds the definition. Nothing is sent from here but an optional correlation id: no workflow, task, provider, mode or digest (FD-10.3). The worker's own definition digest binds the run; a repeated start with the same correlation id replays the same instance.

Correlation id (optional)

Start worker shadow run

Synthetic demonstration data - deterministic planning only - no agent execution - no permission granting - no business-action authorization

Screen 20, as captured 2026-09-07

- **Shows:** two labelled paths: a compiled workflow run in the studio against fixtures, or the governed runtime worker's own shadow run started by relay. Each labelled with its executor, hook and maturity. A permissive hook renders as a red alert saying the result is not a governance result.
- **Answered by:** `v2_simulate_run`; `v2_review_start_shadow_run` relayed to the worker; the agent runtime over the one pinned no-network provider.
- **Operator can:** run a simulation in `DRY_RUN`, `SIMULATION` or `SHADOW`; ask the worker to start its own shadow run.
- **Never:** reaches anything consequential. `LIVE` is absent, not disabled. No path is chosen for the operator.
- **Ruling:** front-door seam 3, FD-7; seam 6, FD-10.

Enter, press, expect

On the live deployment: both paths answer typed gaps: path A `simulation_providers` (no fixture provider registry is configured), path B `review_service` (no governed review service base URL is configured) `[I]`.

STEP	DO THIS	EXPECT
1	Path A: leave Execution mode at <code>DRY_RUN</code> and press Run simulation .	Where a fixture provider is composed: <code>execution_mode DRY_RUN</code> , <code>runtime_id studio-simulation:DRY_RUN</code> , an instance id, <code>governance_hook_configured false</code> , and a trace in which the one consequential task <code>t1</code> BLOCKs with <code>GOVERNANCE_NOT_CONFIGURED</code> . The runtime's default hook blocks; this is the fail-closed default, not a governed run.
2	Path A: repeat with <code>SIMULATION</code> and <code>SHADOW</code> .	Same shape; the mode is threaded into <code>runtime_id</code> and every task argument. <code>LIVE</code> is not in the list and the API refuses it with 422 <code>invalid_execution_mode</code> .
3	Path B: type Correlation id <code>demo-run-1</code> and press Start worker shadow run .	Where the review service is composed: <code>STARTED</code> – parked; a consequential task waits on the review queue, <code>instance_id shadow-<24 hex></code> , <code>workflow_id wf-shadow</code> , <code>definition_digest shadow-v1</code> . Pressing again with the same id: <code>REPLAYED</code> – this correlation id already names an instance.
4	Path B: type <code>bad id!</code> .	Inline: <i>A correlation id is a typed token: letters, digits, '.', '_', ':' and '-', at most 64 characters. Nothing is sent until it is.</i>

Values you can type

- Path A runs the hard-coded workflow `studio-simulation` with one task `t1` (`prepare` on provider `fixture`, consequential) for at most 8 quanta.
- Correlation id pattern `^[A-Za-z0-9][A-Za-z0-9._:-]{0,63}$`; blank is allowed and the worker mints one. The instance id path B returns is what screens 23 and 24 need.

Source: `SimulateScreen.tsx:35-186`; `studio_v2.py:498-581, 747-789`; `deployment/governed-runtime-worker/.../workload.py:110-131`; `backend/tests/test_v2_routes.py:176-289`.

21 Publish (/studio/publish)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data Deterministic planning only No live agent execution No permission granting No business-action authorization Not pilot validated Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Publish

Hand a compiled release package to the console's shadow governed loop and see what came back.

Shadow only. This screen cannot authorize an action, clear one, or run a live loop.

Shadow governed loop

A shadow run observes; it changes nothing in any target system. The studio's console client is restricted to four read and shadow routes, so no other console operation is reachable from this screen even by mistake.

Not available in this deployment: compiled_release

No compiled release exists in this session. Compile a reviewed pack with its approval record on the Policy screen first; only that output can be sent to the shadow loop.

Send to shadow loop

Synthetic demonstration data - deterministic planning only - no agent execution - no permission granting - no business-action authorization

Screen 21, as captured 2026-09-07

- **Shows:** what the console's shadow governed loop returned for a compiled release package.
- **Answered by:** `v2_publish_shadow`; the console over four allowlisted routes, the mode word pinned to `shadow`.
- **Operator can:** send to the shadow loop.
- **Never:** authorizes an action, clears one or runs a live loop. With no console base URL composed, the screen reports the gap.
- **Ruling:** FD-8.1, FD-8.2; CP-3.

Enter, press, expect

On the live deployment: the button is disabled until screen 18 has compiled; once pressed it returns the typed refusal `publish_payload_unmapped` on every deployment, and the console seam is unconfigured `[I]`.

STEP	DO THIS	EXPECT
1	Open the screen before compiling on screen 18.	Local gap <code>compiled_release</code> : <i>No compiled release exists in this session. Compile a reviewed pack with its approval record on the Policy screen first.</i>

STEP	DO THIS	EXPECT
2	Compile on screen 18, return, press Send to shadow loop .	Refused <code>publish_payload_unmapped</code> : <i>a compiled release package is not a governed-loop request: the console's shadow loop needs an assertion, an action and operational signals that a compiled package does not carry, and the studio invents none of them; name a frozen console scenario_id instead.</i>
3	API only: <code>POST /api/v2/publish/shadow</code> with <code>{"scenario_id": "k8s_delete_during_freeze"}</code> .	With a console base URL composed: <code>mode SHADOW</code> and the console's loop result. Without one (the live state): <code>gap console_api</code> <i>no ugence_console_api base URL is configured.</i>

Values you can type

- The screen has no scenario field; the frozen console ids are `k8s_rollout_restart_clean`, `k8s_delete_during_freeze`, `k8s_unsupported_claim` (pattern `^[A-Za-z0-9][A-Za-z0-9._-]{0,63}$`). Run them on the console itself (screen 26).

Source: `PublishScreen.tsx:20-72`; `studio_v2.py:589-642`; `backend/tests/test_publish_pin.py:167-178`.

22 Observe (/studio/observe)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

⚠ Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed ✖

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 · additive contract alongside the frozen v1 explorer · planning, preflight and observation only · no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Observe

Two distinct sources: the governed runtime worker's own audit ledger, and the console's audit chain. Each is labelled with its record type, its executor and its maturity.
This screen re-derives nothing. What is shown is what the worker or the console returned, verification included; no source is chosen for you.

Source A · Worker audit ledger

Record type control-plane audit-ledger rows: receipts and references (kind, payload, digests), raw and uninterpreted; not stage narratives
Executor the governed runtime worker reads its own tenant's ledger (seam 7); the studio names no tenant and re-derives nothing
Maturity REFERENCE_GRADE; durable, per-tenant, hash-chained; tamper-evident, not tamper-proof

Correlation id

Read ledger

Source B · Console audit chain

Record type the console's stage chain: assertion, action, signals and disposition per correlation id
Executor one console instance's in-memory audit store; absent by ruling FD-8.1 until the console is packaged (FD-8.3)
Maturity prototype; lost on restart; never durable, complete, distributed or restart-safe (FD-8.5)

Correlation ids

Not available in this deployment: console_api
no ugence_console_api base URL is configured

Correlation id

Reconstruct

Synthetic demonstration data · deterministic planning only · no agent execution · no permission granting · no business-action authorization

Screen 22, as captured 2026-09-07

- **Shows:** two distinct sources, never merged: the worker's own audit ledger by correlation id, with the worker's chain verification; and the console's audit chain. Each labelled with record type, executor and maturity. Unreachable, empty, not-found and refused are shown differently.
- **Answered by:** v2_observe_ledger_chain relayed to the worker; v2_observe_audit_ids , v2_observe_audit_chain relayed to the console.
- **Operator can:** read a ledger or a chain by correlation id.
- **Never:** re-derives, re-orders or re-hashes anything; chooses no source for the operator.
- **Ruling:** front-door seam 7, FD-11.4 TWO_LABELLED_SOURCES .

Enter, press, expect

On the live deployment: Source A answers the typed gap review_service ; Source B answers the gap console_api (also true of the captured screen) [I].

STEP	DO THIS	EXPECT
1	Source A: type the correlation id you gave screen 20 path B (demo-run-1) and press Read ledger .	Where the worker relay is composed: <code>READ</code> , <code>chain_verified true</code> , <code>entry_count</code> , and rows of <code>entry_ref</code> , <code>kind</code> , <code>recorded_at</code> , <code>recorded_by</code> , <code>record_digest</code> ; labels <i>durable</i> , <i>per-tenant</i> , <i>hash-chained...</i> <i>tamper-evident</i> , <i>not tamper-proof</i> .
2	Source A: type <code>never-ran-1</code> .	<i>The worker reports no entry of its tenant carrying never-ran-1. That is a typed not-found from a reachable worker, not an empty ledger and not a failure to read it.</i>
3	Source B: click a listed correlation id or type one and press Reconstruct .	Where a console is composed: the console's stage chain for that id. Live: the <code>console_api</code> gap in both the id list and the chain.

Values you can type

- Source A ids follow the same typed-token pattern as screen 20; Source B accepts free text.
- Worker refusals pass through verbatim: `REFUSED_INTEGRITY`, `REFUSED_SCHEMA`, `REFUSED_UNCONFIGURED`.

Source: `ObserveScreen.tsx:30-203`; `studio_v2.py:659-673`, `797-841`; `deployment/governed-runtime-worker/.../service.py:428-449`.

23 Review Queue (/studio/review)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.ap1.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data Deterministic planning only No live agent execution No permission granting No business-action authorization Not pilot validated Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.ap1.v2 - additive contract alongside the frozen v1 explorer - planning, preflight and observation only - no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Review Queue

Parked ESCALATE instances awaiting a human decision, as the governed review service lists them.

This screen displays and relays. It holds no approver identity, computes no eligibility, consumes no approval, signals nothing and resumes nothing.

Awaiting a decision

Approver identity: PRESENTED_UNPROVEN

The approver on a decision is a reference the review service listed as eligible and this screen relayed. No identity provider exists, so nothing here proves who decided; the review service records the decision as presented, not proven.

Instance	Task	Operation	Fingerprint (history)	Disposition	Approval	Role	Requested	Expires	Eligible
shadow-1643d255e05d9dda511629cb	t1	shadow-recorder.do	6c401b53d78d...e95b1d	ESCALATE	REQUESTED	risk-approver	2026-09-07T05:05:57.156545+00:00	2026-09-14T05:05:57.156545+00:00	0

Decide

Synthetic demonstration data - deterministic planning only - no agent execution - no permission granting - no business-action authorization

Screen 23, as captured 2026-09-07

- **Shows:** parked ESCALATE instances awaiting a human decision, as the governed review service lists them.
- **Answered by:** `v2_review_list_queue` , `v2_review_submit_decision` , relayed to the review service over seven allowlisted routes.
- **Operator can:** relay a GRANT or REJECT decision with a justification and an opaque proof header, forwarded unread.
- **Never:** holds an approver identity, computes eligibility, consumes an approval, signals or resumes anything. The approver is presented and unproven until an identity provider is validated.
- **Ruling:** HR-1 `DISPLAY_AND_TRANSMIT` ; ID-1 `PASS_THROUGH_OPAQUE_TOKEN` .

Enter, press, expect

On the live deployment: typed gap `review_service` `[I]`. Even where the review service is composed, the reference directory holds no eligible approver, so no decision can be relayed from the screen.

STEP	DO THIS	EXPECT
1	Start a worker shadow run on screen 20 path B, then open the queue.	Where composed: one row: instance <code>shadow-...</code> , task <code>t1</code> , operation <code>shadow-recorder.do</code> , disposition <code>ESCALATE</code> , approval <code>REQUESTED</code> , role <code>approver</code> , requested and expires timestamps, eligible <code>0</code> ; banner <i>Approver identity: PRESENTED_UNPROVEN</i> .
2	Press Decide .	The decision form: Presented approver (a select over the eligible approvers the service listed; with none listed the notice reads <i>The review service reports no eligible approver for approver. Nothing can be relayed until the authority directory reports one.</i>), Decision GRANT or REJECT, Justification (required), Approver proof (optional password field, forwarded once as <code>X-Ugence-Approver-Proof</code> , never stored).
3	With an eligible approver listed: choose it, GRANT, justification <code>Shadow run reviewed for the demo</code> , press submit.	<i>Recorded by the review service: RECORDED</i> , <code>signal_delivered true</code> , <code>resume_delivered true</code> ; the instance leaves the queue and screen 24 shows the decision event.
4	Submit the identical decision again.	<code>REPLAYED</code> . A different decision on the same approval: <code>REFUSED_ALREADY_DECIDED</code> (<i>the first decision stands</i>).

Values you can type

- Only path B on screen 20 produces a queue item; path A blocks in-process and never parks.
- Required role in the reference deployment: `approver` (`UGENCE_REVIEW_REQUIRED_ROLE`). The scope that must hold it is not fixed by any repository value, and the directory has no seed path, so eligibility stays empty until an administrator loads a grant behind the worker's identity gate (AP-3).

Refusals to expect

- `REFUSED_INELIGIBLE`, `REFUSED_UNKNOWN_APPROVAL`, `REFUSED_NOT_REVIEWABLE`, `REFUSED_NOT_OPEN`, `REFUSED_ALREADY_DECIDED`, `REFUSED_INVALID_DECISION`; with an identity port composed also `REFUSED_UNAUTHENTICATED`, `REFUSED_IDENTITY_MISMATCH`, `REFUSED_NOT_HUMAN`, `REFUSED_TENANT_UNPROVEN`.

Source: `ReviewQueueScreen.tsx:36-332`; `studio_v2.py:710-721`; `packages/integration/governed-review-service/.../service.py:238-251, 487-676`; `RAILWAY_REFERENCE_DEPLOYMENT.md` part 7.7.

24 Run Detail (/studio/review/:instanceId)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data · Deterministic planning only · No live agent execution · No permission granting · No business-action authorization · Not pilot validated · Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 · additive contract alongside the frozen v1 explorer · planning, preflight and observation only · no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Run Detail

Instance shadow-1643d255e05d9dda511629cb, as the governed review service renders it.

Reads only. This screen resumes, releases, continues and signals nothing; fingerprints and valid_until values are history, never a live permission.

← Review Queue

Instance

Shown as history. A fingerprint is what was evaluated and a valid_until is when that evaluation lapsed. Neither is a live permission: a pre-park clearance is never reused, and the next evaluation is fresh.

Parked on an ESCALATE. A human decision is awaited; see the open approvals below and the Review Queue.

Workflowwf-shadow

StatusPAUSED

Correlation idc-explainer-1

Engine{"known":true,"engine_id":"dbos","workflow_id":"wf-shadow","definition_digest":"shadow-v1","state_writes":2,"event_count":9}

Identity proofPRESENTED_UNPROVEN

Task	Status	Attempts	Disposition	Operation	Fingerprint (history)	valid_until (history)	Evaluation ref
t1	WAITING	0	ESCALATE	shadow-recorder.do	6c401b53d78d_e95b1d	—	—

Open approvals

apr_07d70fbba7970a3220ddd77e0111b33d · REQUESTED

Receipt linkages

Shown as history. A linkage joins the proposal fingerprint, the approval, its consumption, the park, the decision signal, the resume and the two evaluations; the audit reference says where the review service wrote it. This screen writes nothing.

The review service reports no decided approval for this instance, so there is no linkage to show.

Runtime events

Seq	Event	Attempt	Body
1	WORKFLOW_CREATED	—	{"seq":0,"type":"WORKFLOW_CREATED","detail":{"instance_id":"shadow-1643d255e05d9dda511629cb","workflow_id":"wf-shadow"}}

Screen 24, as captured 2026-09-07

- **Shows:** one run, its events, its approval, and the HE-1 linkage as the review service exposes them.
- **Answered by:** v2_review_read_run , v2_review_read_run_events , v2_review_read_approval .
- **Operator can:** read.
- **Never:** resumes, releases, continues or signals; fingerprints and valid-until values are history, never a live permission.
- **Ruling:** GAS-7 HR-D; HE-5.

Enter, press, expect

On the live deployment: typed gap review_service [I].

Ugence Developer Agentic Studio Explainer, with inputs · 2026-09-13

53 / 78

25 Status (/studio/status)

Ugence Governance Studio

Eligibility Explorer

Governed Agent Studio

governance_studio.api.v1

Eligibility

Selected

Assigned

Authorized

Executed

Synthetic demonstration data Deterministic planning only No live agent execution No permission granting No business-action authorization Not pilot validated Not production certified

Registration

Data use

Vendor dependencies

Constitution

Policy

Authority

Simulate

Publish

Observe

Review

Status

governance_studio.api.v2 · additive contract alongside the frozen v1 explorer · planning, preflight and observation only · no screen here issues, activates, revokes, grants, authorizes, clears or executes; the review screens display and relay, and signal or resume nothing

Status

What this deployment attested about itself before its port bound: the six front-door seam states, the integrity checks and the pinned identities, handed to the studio once at composition. This panel probes nothing, changes nothing, and lists no module registry. A configured seam is not a reachable engine, and nothing here is read at request time.

What this is evidence of

STARTUP_ATTESTATION: what this deployment attested about itself before its port bound, handed to the studio once at composition and never re-read. A configured seam is not a reachable engine, a passed check is not a live one, and nothing here is probed at request time.

Gate result

PASS · OK

Seam states, as the gate attested them

Constitution registry (seam 1) constitution_registry

Authority reads (seam 2) authority_reads

Simulation provider (seam 3) simulation_provider

System registry (seam 5) system_registry

Data-use declarations (seam 8) data_use_declarations

Vendor declarations (seam 9) vendor_declarations

"configured" means the seam's file or flag was present and writable when the gate ran. It is not a statement that any engine is reachable now.

configured

unset

configured

unset

unset

unset

Integrity checks

passed config_valid

passed credentials_configured

passed allowed_hosts_configured

passed no_dev_mode_in_production

passed constitution_registry_writable

passed simulation_no_permissive_hook

passed frontend_build_exists

passed frontend_version_0_2_0

passed backend_api_version_0_1_0

passed api_contract_governance_studio_api_v1

passed openapi_hash_unchanged

passed approved_operation_manifest_valid

Screen 25, as captured 2026-09-07

- **Shows:** what the deployment attested about itself before its port bound: the six front-door seam states, the integrity checks, the pinned identities, under a ceiling stating that a configured seam is not a reachable engine.
- **Answered by:** `v2_observe_deployment` ; the integrity gate's report, handed to the studio once at composition.
- **Operator can:** read.
- **Never:** probes anything, changes anything, or lists a module registry.
- **Ruling:** MA-2 as amended by MS-1 to MS-5; v2-A7.

Enter, press, expect

On the live deployment: typed gap `deployment_report` : no startup integrity report was handed to the studio at composition; this process was not started through the deployment's integrity gate `[I]`.

STEP	DO THIS	EXPECT
1	Open the screen.	Live: the gap above, which itself confirms the bare start command.

STEP	DO THIS	EXPECT
2	Under the Docker/compose profile.	<code>result PASS</code> or <code>FAIL</code> , a <code>checks</code> map, the pinned identities (deployment version, frontend build hash, API contract, OpenAPI sha256, synthetic bundle hash), and six seam chips: <code>constitution_registry</code> , <code>authority_reads</code> , <code>simulation_provider</code> , <code>system_registry</code> , <code>data_use_declarations</code> , <code>vendor_declarations</code> , each <code>configured</code> , <code>unwritable</code> or <code>unset</code> ; ceiling a <i>configured seam is not a reachable engine</i> .

Values you can type


- No inputs. Setting `UGENCE_STUDIO_SIMULATION_PROVIDER=1` alone in the compose profile flips `simulation_provider` to `configured` and turns screen 20 path A on.

Source: `StatusScreen.tsx:24-135`; `studio_v2.py:1505-1592`; `deployment/governance-studio/.../startup_integrity.py:90-243`.

C · Console (3 tabs, the packaged console API)

The console app is a separate deployable. It calls exactly four of the five routes the packaged `ugence-console-api` serves; the six routes ruling CP-3 withheld are not called, and the two it used to call show a typed gap. Every answer carries the service's audit ceiling: one process's view, lost on restart.

26 Governed Loop



Screen 26, as captured 2026-09-07

- **Shows:** the shadow governed loop's stage trail for one scenario: gateway, verify, authorize, clear, record, with each module's decision and the final shadow disposition.
- **Answered by:** `POST /v1/governed-loop/scenario/{scenario_id}`.
- **Operator can:** type a scenario id and run the loop in shadow. The catalogue is a typed gap; an unknown id is the service's own 404.
- **Never:** executes anything; the loop evaluates and records, and the mode is shadow.
- **Ruling:** CP-3; MA-4 `RETIRE_THE_TWO_CALLS`.

Enter, press, expect

On the live deployment: works in full `[V]` per the deployment reference and the demo plan.

STEP	DO THIS	EXPECT
1	Open <code>https://console-web-production-4e77.up.railway.app</code> , tab Governed Loop . Read the amber <code>SHADOW MODE</code> pill and the <code>CONSOLE_ROUTE_WITHHELD</code> panel (the catalogue route <code>GET /v1/scenarios</code> is deliberately not served).	The scenario id box is empty; Run governed loop is disabled until it has text.


STEP	DO THIS	EXPECT
2	Type <code>k8s_rollout_restart_clean</code> ; press Run governed loop (or Enter).	Banner <i>OBSERVED (shadow)</i> — enforcement would <i>ALLOW</i> ; nothing changed. with correlation <code>corr-<12hex></code> and <code>cer-942ee6db6bbb5cb9</code> . Stages: Gateway <i>ADMITTED Admitted 3/4 units (1 redundant dropped, lossless=True)</i> ; Verify <i>SUPPORTED Assertion supported at evidence coverage 100%</i> ; Authorize <i>AUTHORIZED reasons ['policy_allow']</i> ; Clear <i>CLEAR</i> ; <i>OPERATIONALLY_SAFE</i> ; Record <i>RECORDED</i> .
3	Type <code>k8s_delete_during_freeze</code> ; run.	enforcement would <i>BLOCK</i> ; <code>cer-9d5c7ca5e6c21802</code> . Gateway, Verify and Authorize identical to step 2, including <i>AUTHORIZED ... policy_allow</i> ; Clear <i>HOLD</i> ; <i>Operational clearance HOLD: CHANGE_FREEZE_ACTIVE, ERROR_BUDGET_EXHAUSTED</i> . Policy said yes; clearance said no.
4	Type <code>k8s_unsupported_claim</code> ; run.	enforcement would <i>BLOCK</i> ; <code>cer-202977ea757e69c9</code> . Verify <i>INDETERMINATE Assertion indeterminate at evidence coverage 0% (no evidence refs)</i> ; Authorize still <i>AUTHORIZED ... policy_allow</i> ; Clear <i>CLEAR</i> . Policy said yes; the evidence gate said the claim is not supported. The word on screen is <i>INDETERMINATE</i> , not <i>UNSUPPORTED</i> .
5	Type <code>k8s_foo</code> ; run.	Red box, verbatim: <code>Error: 404 {"detail":"unknown scenario 'k8s_foo'"}.</code>

Values you can type

- Exactly three scenario ids exist: `k8s_rollout_restart_clean` (*ALLOW*), `k8s_delete_during_freeze` (*HOLD* → *block*), `k8s_unsupported_claim` (*INDETERMINATE* → *block*). CER ids are stable per scenario; correlation ids are fresh per run.
- Copy the correlation id from the banner for screen 28.

Source: `apps/console/src/views/GovernedLoop.tsx:32-146`; `packages/integration/console-api/.../scenarios.py:33-117`; `orchestrator.py:56-141`; `capabilities/operational_safety.py:20-63`; `docs/deployment/CLIENT_DEMO_PLAN.md:23-42`.

27 Modules



Screen 27, as captured 2026-09-07

- **Shows:** the service's availability probes for the four engines it can reach, under the keys it returns, and its declared audit ceiling. The nine-module registry is a typed gap.
- **Answered by:** `GET /health`.
- **Operator can:** read.
- **Never:** lists module descriptions, maturity or wiring; those live behind a withheld route.
- **Ruling:** CP-3, CP-4; MA-4.

Enter, press, expect

On the live deployment: works in full `[V]`.


STEP	DO THIS	EXPECT
1	Open tab Modules .	<code>CONSOLE_ROUTE_WITHHELD</code> panel for <code>GET /v1/modules</code> , then <i>Availability probes · from GET /health</i> : four cards, in order <code>context_minimization</code> , <code>tap</code> , <code>actiongate</code> , <code>autonomous_control_plane</code> , each with a green dot and <i>engine available</i> .

STEP	DO THIS	EXPECT
2	Read the ceiling line.	<i>Audit ceiling, as the service declares it on every answer: IN_MEMORY_SINGLE_PROCESS: this audit is one process's view of its own runs and is lost on restart. It is not a durable, tamper-evident or shared record.</i>

Values you can type

- No inputs. A grey dot with an import error string means that engine's package did not install;
`autonomous_control_plane` is always available.

Source: `apps/console/src/views/Modules.tsx:17-63`; `console-api/.../app.py:103-115`; `models.py:23-26`.



Screen 28, as captured 2026-09-07

- **Shows:** a decision chain reconstructed by correlation id: what was asserted, whether it was supported, who authorized, whether it was safe.
- **Answered by:** GET /v1/audit , GET /v1/audit/{correlation_id} .
- **Operator can:** pick or type a correlation id.
- **Never:** re-derives a chain; the console's in-memory store is the record.
- **Ruling:** CP-4 DECLARE_THE_CEILING .

Enter, press, expect

On the live deployment: works in full [V], but only for runs made in the current service process.

STEP	DO THIS	EXPECT
1	Run at least one scenario on screen 26 first, then open tab Audit.	A row of red chips, one per correlation id from this process, newest last. Before any run the chip row is absent.

STEP	DO THIS	EXPECT
2	Click a chip, or paste a <code>corr-...</code> id and press Reconstruct .	<code>corr-...</code> · <code>cer-...</code> · <code>mode shadow</code> , the final disposition line, and a four-row table Stage / Module / Decision / Summary: Gateway, Verify, Authorize, Clear. The Record stage is not in the stored chain, so the audit shows four rows where the loop showed five.
3	Paste <code>corr-deadbeef</code> and press Reconstruct .	Error: 404 {"detail":"no record for 'corr-deadbeef'"}.
4	Press Reconstruct with the box empty.	Error: 404 {"detail":"Not Found"} (the button is not disabled when empty).

Values you can type

- No ids are pre-seeded; a service restart empties the store (CP-4).

Source: `apps/console/src/views/Audit.tsx:8-102`; `console-api/.../audit.py:28-43`; `orchestrator.py:111-125`; `app.py:143-148`.

D · Authority Plane (4 screens, the governed runtime worker's reads)

The administrators' surface, its own deployable (AP-2). At this step it reads and does not write. Every answer is shown under an identity banner repeating the worker's own words: the read was not authenticated; the decision proof the deployment can give; the adapter's issuer validation, still in-process only; and that a grant is what an administrator loaded. The worker implements loading and revoking a role grant behind its identity gate (AW-2 to AW-5) but serves neither until the adapter is validated against a real enterprise issuer (AP-3 controlling, §18); this app has no write control and its client cannot name a write.

29 Grants

Ugence Authority Plane

Administrators' surface · reads only until enterprise issuer validation · over the governed runtime worker

AP-5 READS_FIRST · AP-3 IDP_VALIDATED_FIRST

Grants

Holders

Committee

Grant events

Reads only.

Loading a role grant and revoking one are implemented on the worker behind its identity gate, but are not served until the identity adapter has been validated end to end against a real enterprise issuer (AP-3); the worker's in-process issuer evidence is implementation and conformance evidence only. Activating a constitution and issuing a record act on stores the worker does not compose (AW-2). Nothing on this surface authorizes, clears or executes anything, at any step (AP-4).

Grants held by a principal

The role grants one principal holds in this worker's tenant, active at the worker's clock. The principal id is the issuer-qualified subject the identity adapter presents, percent-encoded.

https%3A%2F%2Fidp.example%7Calice

Read grants

read_authenticated: false

decision proof: PRESENTED_UNPROVEN

issuer validation: IN_PROCESS_ISSUER_ONLY

REFERENCE_GRADE_SHADOW_ONLY

This read was not authenticated. No identity port is composed, so every decision on this worker is a presented, unproven approval; the adapter has been validated against an in-process issuer only. What is listed below is what an administrator loaded; the directory attests nothing about whether it should exist.

tenant tenant-a · as of 2026-09-07T06:24:17.267410+00:00

https%3A%2F%2Fidp.example%7Calice holds 1 active grant.

PRINCIPAL	KIND	ROLE	SCOPE	ISSUED	EXPIRES	LOADED BY	GRANT ID
https%3A%2F%2Fidp.example%7Calice	HUMAN	risk-approver	approval/policy_pack	2026-09-04T05:04:56.487268+00:00	2026-12-06T05:04:56.487268+00:00	admin-load	grant_c0241af4cd21de04

▶ the worker's answer, verbatim

REFERENCE_GRADE_SHADOW_ONLY · a grant shown here is what an administrator loaded · no identity provider is provisioned · no write is served before enterprise issuer validation

Screen 29, as captured 2026-09-07

- **Shows:** the role grants one principal holds in the worker's tenant, active at the worker's clock.
- **Answered by:** GET /authority/grants?principal_id=; the sqlite authority directory the worker composed.
- **Operator can:** type a principal id and read; follow a grant to its events.
- **Never:** loads or revokes a grant. Those writes are implemented on the worker and not served before enterprise issuer validation (AP-3).
- **Ruling:** AP-5 READS_FIRST.

Enter, press, expect

On the live deployment: works [V]; the directory is empty, so every principal holds zero grants. The captured screen shows a seeded local directory (tenant-a , principal https%3A%2F%2Fidp.example%7Calice , role risk-approver) that does not exist on the live worker.

Ugence Developer Agentic Studio Explainer, with inputs · 2026-09-13

65 / 78

STEP	DO THIS	EXPECT
1	Open <code>https://authority-plane-production.up.railway.app</code> , tab Grants ; type <code>approver</code> and press Read grants .	HTTP 200. Identity banner: chips <code>read_authenticated: false</code> , decision proof: <code>PRESENTED_UNPROVEN</code> , issuer validation: <code>IN_PROCESS_ISSUER_ONLY</code> , <code>REFERENCE_GRADE_SHADOW_ONLY</code> ; line <code>tenant tenant-demo · as of <instant></code> . Then <i>approver holds 0 active grants. and No grant of this tenant matches, at this instant. An empty list is what the directory holds, not a refusal. Expand the worker's answer, verbatim for the JSON (result READ, grant_count 0).</i>
2	Type <code>https%3A%2F%2Fidp.example%7Calice</code> and read.	Same: zero grants. The value from the capture is test-only.
3	Leave the box empty.	The button is disabled; nothing is sent.
4	Browser console on the plane's origin: <code>fetch('/api/authority/grants', {method: 'POST'}).then(r=>console.log(r.status))</code> .	<code>405</code> ; body <code>{"error":"method_not_allowed","detail":"this front forwards reads only"}</code> . There is no write path.

Values you can type

- A principal id is any typed token: non-empty, at most 256 characters, no whitespace, NFC. The browser percent-encodes it, so a space becomes `%20` and is still a token (200, zero rows); the 422 `REFUSED_UNTYPED` is reachable only with curl.
- With a grant present the table shows Principal, Kind, Role, Scope, Issued, Expires, Loaded by, Grant id and an **Events** button.

Source: `apps/authority-plane/src/views/GrantsView.tsx:7-51`; `components/Identity.tsx:12-51`; `server.mjs:120-165`; `deployment/governed-runtime-worker/.../authority_reads.py:56-152`; `RAILWAY_REFERENCE_DEPLOYMENT.md` part 7.7.

30 Holders

Ugence Authority Plane

Administrators' surface · reads only at this step · over the governed runtime worker

AP-5 READS_FIRST

Grants

Holders

Committee

Grant events

Reads only.

Loading a grant, revoking one, activating a constitution and issuing a record are the plane's writes, and they are not served until the identity adapter is validated against a real issuer and every write carries an issuer-authenticated subject (AP-3). Nothing on this surface authorizes, clears or executes anything, at any step (AP-4).

Holders of a role in a scope

Every principal of this tenant holding one role in one scope at the worker's clock, committees included.

risk-approver

approval/policy_pack

Read holders

read_authenticated: false

decision proof: PRESENTED_UNPROVEN

issuer validation: IN_PROCESS_ISSUER_ONLY

REFERENCE_GRADE_SHADOW_ONLY

This read was not authenticated. No identity port is composed, so every decision on this worker is a presented, unproven approver; the adapter has been validated against an in-process issuer only. What is listed below is what an administrator loaded; the directory attests nothing about whether it should exist.

tenant tenant-a · as of 2026-09-07T05:09:30.751346+00:00

3 holders of risk-approver in approval/policy_pack.

PRINCIPAL	KIND	ROLE	SCOPE	ISSUED	EXPIRES	LOADED BY	GRANT ID
https%3A%2F%2Fidp.example%7Calice	HUMAN	risk-approver	approval/policy_pack	2026-09-04T05:04:56.487268+00:00	2026-12-06T05:04:56.487268+00:00	admin-load	grant_c0241af4cd21c
https%3A%2F%2Fidp.example%7Cbob	HUMAN	risk-approver	approval/policy_pack	2026-09-04T05:04:56.487268+00:00	2026-12-06T05:04:56.487268+00:00	admin-load	grant_c8579d5e3437i
risk-committee	COMMITTEE	risk-approver	approval/policy_pack	2026-09-04T05:04:56.487268+00:00	2026-12-06T05:04:56.487268+00:00	admin-load	grant_9a6a358ce1ccf

► the worker's answer, verbatim

REFERENCE_GRADE_SHADOW_ONLY · a grant shown here is what an administrator loaded · no identity provider is provisioned

Screen 30, as captured 2026-09-07

- **Shows:** every principal of the tenant holding one role in one scope, committees included.
- **Answered by:** GET /authority/holders?role=&scope=.
- **Operator can:** type a role and scope and read.
- **Never:** changes a holder.
- **Ruling:** AP-5.

Enter, press, expect

On the live deployment: works [V]; zero holders on the empty directory.

STEP	DO THIS	EXPECT
1	Tab Holders: role <code>approver</code> , scope <code>approval/policy_pack</code> ; press Read holders (Enter in the scope field also submits).	HTTP 200, identity banner, <i>0 holders of approver in approval/policy_pack.</i> , the empty list ... not a refusal notice, verbatim JSON with <code>holder_count 0</code> .
2	Leave either field blank.	The button stays disabled.

Values you can type

- `approver` is the deployment's own required role (`UGENCE_REVIEW_REQUIRED_ROLE`). No scope value is fixed by the repository; any typed token returns zero on the live directory.

Source: `HoldersView.tsx:11-47`; `authority_reads.py:156-164`.

31 Committee

Ugence Authority Plane

Administrators' surface · reads only at this step · over the governed runtime worker

AP-5 READS_FIRST

Grants

Holders

Committee

Grant events

Reads only.

Loading a grant, revoking one, activating a constitution and issuing a record are the plane's writes, and they are not served until the identity adapter is validated against a real issuer and every write carries an issuer-authenticated subject (AP-3). Nothing on this surface authorizes, clears or executes anything, at any step (AP-4).

Committee report

One committee's members holding one role in one scope, counted against the committee's quorum at the worker's clock. A member whose grant lapsed or was revoked is not counted.

risk-committee

risk-approver

approval/policy_pack

Read committee

read_authenticated: false

decision proof: PRESENTED_UNPROVEN

issuer validation: IN_PROCESS_ISSUER_ONLY

REFERENCE_GRADE_SHADOW_ONLY

This read was not authenticated. No identity port is composed, so every decision on this worker is a presented, unproven approver; the adapter has been validated against an in-process issuer only. What is listed below is what an administrator loaded; the directory attests nothing about whether it should exist.

tenant tenant-a · as of 2026-09-07T05:09:32.734189+00:00

risk-committee · quorum 2 · members counted 1 · quorum not met at this instant

PRINCIPAL	KIND	ROLE	SCOPE	ISSUED	EXPIRES	LOADED BY	GRANT ID
https%3A%2F%2Fidp.example%7Cbob	HUMAN	risk-approver	approval/policy_pack	2026-09-04T05:04:56.487268+00:00	2026-12-06T05:04:56.487268+00:00	admin-load	grant_c8579d5e3437ac4f8b1

▶ the worker's answer, verbatim

REFERENCE_GRADE_SHADOW_ONLY · a grant shown here is what an administrator loaded · no identity provider is provisioned

Screen 31, as captured 2026-09-07

- **Shows:** one committee's counted members against its quorum, at this instant; a lapsed or revoked member is not counted.
- **Answered by:** GET /authority/committees/{committee_id}?role=&scope=.
- **Operator can:** read; a missing committee is a typed not-found.
- **Never:** adds a member or changes a quorum.
- **Ruling:** AP-5; directory ruling D-4.

Enter, press, expect

On the live deployment: works [V]; answers a typed 404 on the empty directory, which is the correct result.

STEP	DO THIS	EXPECT
1	Tab Committee : committee id risk-committee, role approver, scope approval/policy_pack; press Read committee .	HTTP 404. No identity banner; instead the grey notice NOT_FOUND · HTTP 404 · no committee 'risk-committee' holds a grant of 'approver' in 'approval/policy_pack' for this tenant at <instant>.


STEP	DO THIS	EXPECT
2	Leave any field blank.	The button stays disabled.

Values you can type

- With a committee loaded the line reads `<committee> · quorum N · members counted M · quorum met` (or *not met at this instant*); revoked or lapsed members are not counted.

Source: `CommitteeView.tsx:12-55`; `authority_reads.py:91, 174-179`; `CLIENT_DEMO_PLAN.md:122-128`.

32 Grant events



Screen 32, as captured 2026-09-07

- **Shows:** a grant's append-only history: loaded, and possibly revoked, with the actor recorded. A grant of another tenant reads as unknown.
- **Answered by:** `GET /authority/grants/{grant_id}/events`.
- **Operator can:** read.
- **Never:** edits a grant in place; the directory records loading and revocation only.
- **Ruling:** AP-5; AP-3 for the writes it does not have.

Enter, press, expect

On the live deployment: works `[V]`; answers a typed 404 for every grant id on the empty directory.

STEP	DO THIS	EXPECT
1	Tab Grant events : type <code>grant_abc123</code> ; press Read events .	HTTP 404, grey notice <code>NOT_FOUND</code> · HTTP 404 · this tenant holds no grant <code>'grant_abc123'</code> .
2	Arrive via an Events button on a grant row (only possible with data).	The field is pre-filled and read once: the grant row, then an ordered list <code><seq></code> · <code>GRANTED</code> · <code><instant></code> · by <code><actor></code> · <code><detail></code> , later <code>REVOKED</code> ; append-only.

Values you can type

- A grant id of another tenant returns the same not-found, deliberately indistinguishable from an unknown id.

Source: `EventsView.tsx:12-61`; `authority_reads.py:189-192`; `deployment/governed-runtime-worker/tests/test_authority_reads.py:133-146`.

E · Governance Studio, Bring Your Workflow (1 screen, contract v1)

The one screen on which a document the operator supplies enters the studio. Owner ruling BW-1 to BW-5 (§22) admitted it as a read-only Workflow IR inspector and superseded, for this surface only, the earlier "no arbitrary JSON / fixture-upload input" sentences. The frozen v1 OpenAPI document did not change: three operations it already carried moved from the front end's forbidden list to its approved list. Owner ruling BW-3A (§24, 2026-09-10) split phase 3 and shipped its drafts half: the screen may keep a server-validated document as an unapproved DRAFT for the deployment's own tenant through three operations of the v2 contract (amendment v2-A8); the identity half (verified owners, directory grants, submit for approval) is phase 3B, behind AP-3.

33 Bring Your Workflow (/bring-your-workflow)

Ugence Governance Studio

Eligibility Explorer

Bring Your Workflow

Governed Agent Studio

governance_studio.api.v1

🔔 Eligibility ✖ Selected ✖ Assigned ✖ Authorized ✖ Executed ✖

Synthetic demonstration data

Deterministic planning only

No live agent execution

No permission granting

No business-action authorization

Not pilot validated

Not production certified

Bring Your Workflow

REFERENCE_GRADE

ephemeral · no execution

Validate and adapt an existing agentic workflow for Ugence governance.

Accepts Ugence Workflow IR JSON. It does not execute, publish or persist the submitted workflow.

Load the guided example

Clear

The guided example is the procurement compiled workflow the scenario catalog serves, bundled with this screen.

WORKFLOW IR JSON

```
{
  "manifest": {
    "policy_pack_id": "gs_procurement",
    "policy_pack_version": 1,
    "structural_digest": "sha256:179d78b8396339d2650ccad4af2865458c00749990c1a8a5d09b2f4ae62811e8"
  },
  "release_metadata": {
    "synthetic": true
  },
  "structural_digest": "sha256:179d78b8396339d2650ccad4af2865458c00749990c1a8a5d09b2f4ae62811e8",
  "workflow_ir": {

```

or choose a local JSON file

Choose File

 No file chosen

WHAT THE DOCUMENT DECLARES

declared version

workflow_ir.v1

size

6048 bytes · depth 6 · 178 values

nodes / edges

9 nodes · 8 edges

node kinds

ACTION_CLEARANCE_REQUIREMENT × 1

ACTION_CONSTRAINT × 1

APPROVAL_GATE × 1

AUDIT_EMISSION × 1

DECISION_RULE × 1

EVIDENCE_REQUIREMENT × 3

TERMINAL_OUTCOME × 1

declared dispositions

ADVISORY × 6

AUTHORITATIVE × 3

human review required

none declared

human authority required

proc_binding_approval

required tool refs

none declared

referenced capabilities

ACTION_CLEARANCE · ACTION_GATE · COMPTIFR · DECISION_AUTHORITY

Screen 33, as captured 2026-09-07

- **Shows:** what a pasted or locally chosen Ugence Workflow IR JSON document declares (version, size, nodes, edges, node kinds, dispositions, human review and authority requirements, tool and capability refs, policy pack, fingerprint, canonical digest); then, on request, the server's validation, adaptation (adapter mode, node dispositions, role requirements, fingerprints, diagnostics) and comparison of a v1 and a v2 adaptation; and, since BW-3A, the draft the server kept (derived id, lifecycle `DRAFT`, kept digest, record digest, the claimed owner with its assurance `PRESENTED_UNPROVEN`, lineage, what it confers: nothing) and the drafts the deployment keeps, read on request and never on load. The disclaimer, verbatim: "Accepts Ugence Workflow IR JSON. It does not execute, compile, approve or publish the submitted workflow. A validated document may be kept only as an unapproved DRAFT for this deployment's tenant."
- **Answered by:** `validate_workflow`, `adapt_workflow`, `compare_adaptations` on the v1 contract and `v2_workflow_drafts_save`, `v2_workflow_drafts_list`, `v2_workflow_drafts_read` on the v2 contract; nothing else. The guided example is the procurement compiled workflow the catalog serves, bundled with the screen.
- **Operator can:** paste JSON or choose a local `.json` file (read in the browser); load the guided example; validate; adapt; compare against the same workflow in the other contract version; download the report and the adapted envelope to their device; keep the gated document as a draft with a title, a claimed owner, an optional link to an AI-

system registration (reference plus digest) and an optional predecessor; list the deployment's drafts; load a kept draft back through the gate, which pre-fills the revision's lineage.

- **Never:** executes, simulates, compiles, approves, publishes or exports the document; keeps the pasted text (the server keeps its own canonical encoding, validated first); keeps a draft under a tenant the browser names (the tenant is server configuration); authenticates an owner; edits or deletes a draft (a revision supersedes); fetches a URL; reads an archive; accepts code, YAML, a framework-native object or a credential-shaped value; adds to or changes the scenario catalog; converts from LangGraph, CrewAI, AutoGen, n8n or BPMN in the browser. Both the browser gate and the server refuse a document over 1 MiB, deeper than 32 levels, or with more than 200 nodes or 400 edges; the server's refusal is the typed 422 `workflow_too_complex`. Without a drafts file configured, the draft controls report the typed gap `workflow_drafts`.
- **Ruling:** BW-1 `READ_ONLY_WORKFLOW_IR_INSPECTOR`, BW-2 `PASTE_OR_LOCAL_FILE_BODY_ONLY` with the owner's figures, BW-3 `VALIDATE_ADAPT_COMPARE_ONLY`, BW-4 `EPHEMERAL_NO_SERVER_STORAGE` as amended for the one draft write, BW-5 `REFERENCE_GRADE` (\$22); BW-3A `DRAFTS_AUTHORIZED_NOW` with BW-3A.1 to BW-3A.5 (\$24).

Enter, press, expect

On the live deployment: works in full.

STEP	DO THIS	EXPECT
1	Open <code>https://studio-web-production-65e8.up.railway.app/bring-your-workflow</code> ; press Load the guided example .	The procurement compiled workflow fills the textarea. <i>What the document declares:</i> version <code>workflow_ir.v1</code> ; 9 nodes, 8 edges; node kinds <code>ACTION_CLEARANCE_REQUIREMENT</code> ×1, <code>ACTION_CONSTRAINT</code> ×1, <code>APPROVAL_GATE</code> ×1, <code>AUDIT_EMISSION</code> ×1, <code>DECISION_RULE</code> ×1, <code>EVIDENCE_REQUIREMENT</code> ×3, <code>TERMINAL_OUTCOME</code> ×1; dispositions <code>ADVISORY</code> ×6, <code>AUTHORITATIVE</code> ×3; human authority required <code>proc_binding_approval</code> ; capabilities <code>ACTION_CLEARANCE</code> , <code>ACTION_GATE</code> , <code>COMPILER</code> , <code>DECISION_AUTHORITY</code> ; policy pack <code>gs_procurement@1</code> ; fingerprint <code>sha256:179d78b8...11e8</code> .
2	Press Validate .	State <code>VALID</code> ; <i>server supports workflow_ir.v1, workflow_ir.v2; integrity client ... · server sha256:07a01038...d86d · match</i> ; no diagnostics.
3	Press Adapt .	<code>ok yes</code> ; adaptation fingerprint <code>sha256:bfec6412...bf4e</code> , envelope <code>sha256:fa7ea8fe...70b30</code> ; 9 node dispositions with reason codes such as <code>deterministic_kind:DECISION_RULE</code> ; 3 role requirements (<code>role::proc_supplier_evidence</code> , <code>role::proc_supplier_risk</code> , <code>role::proc_recommendation</code>).
4	Paste the v2 counterpart of the same workflow into the second textarea (see values) and press Compare adaptations .	<code>equivalence_state SEMANTICALLY_EQUIVALENT</code> , <code>differences: []</code> , v1 <code>sha256:91c1f7d2...e25a</code> , v2 <code>sha256:ae122481...317f</code> .
5	Press Download report and Download adapted envelope .	Two JSON files saved locally; nothing was stored on the server.

STEP	DO THIS	EXPECT
6	With the gate still passing, fill Keep as an unapproved draft (phase 3A): Title <code>procurement baseline</code> , Claimed owner <code>directory://people/owner-1</code> , leave both registration fields and Supersedes empty, Notes free text. Press Keep as draft .	The one v2 draft write is sent. <i>draft id</i> a derived identifier; <i>lifecycle</i> <code>DRAFT</code> with its note; <i>kept digest</i> and <i>record digest</i> ; <i>integrity</i> client and server agree; <i>claimed owner</i> <code>directory://people/owner-1</code> . <code>PRESENTED_UNPROVEN</code> ; <i>supersedes</i> none, the first of its lineage; <i>registration link</i> none; <i>recorded by</i> and <i>validated by</i> ; <i>confers</i> nothing. The pasted text is never kept — the server keeps its own canonical encoding, validated again first.
7	Press Show kept drafts in <i>Drafts this deployment keeps</i> .	The count, <code>lifecycle DRAFT</code> , the claimed-owner status and what a draft confers; then one row per draft with its id, title, claimed owner, supersedes and superseded-by. Read on request, never on page load, and only this deployment's own tenant is ever answered.
8	Press Load on that row.	The draft returns through the same gate as anything pasted. The form pre-fills and Supersedes is set to the loaded draft id, with the note <i>keeping it again records a revision that supersedes it</i> . Nothing is edited in place.
9	Change Title to <code>procurement baseline rev 2</code> and press Keep as draft again.	A second draft whose <i>supersedes</i> names the first. Tick include superseded revisions and press Show kept drafts to see both; unticked, only the head of the lineage is listed.
10	Press Clear , paste <code>{"ir_version":"workflow_ir.v9"}</code> .	Gate refusal <code>UNSUPPORTED_VERSION: declared version "workflow_ir.v9" is not one of workflow_ir.v1, workflow_ir.v2</code> ; Validate and Adapt disabled.

Values you can type

- Smallest document the browser gate accepts: `{"ir_version":"workflow_ir.v1","nodes":[],"edges":[]}` (the server then reports the adapter's diagnostics).
- A v2 counterpart for Compare: `frontend/tests/fixtures/bring.example-v2.json` or `backend/.../data/conformance_v2/procurement/v2_workflow.json` (top-level `ir_version` `workflow_ir.v2`, `base_ir`, nine `node_semantics`). Compare needs exactly one v1 and one v2 document.
- Draft form: **Title** is the only required field. **Claimed owner** is an opaque handle recorded as `PRESENTED_UNPROVEN`; it confers no read, write, approval or execution authority. An **AI-system registration id** must be accompanied by the **Registration record digest** of that exact record, and both are matched against this tenant's own registry — use the identifier and record digest from screen 14. **Supersedes** must name the head of its lineage.
- The tenant is never sent from the browser on any draft operation; it is this deployment's server configuration.
- Browser gate limits: 1 MiB, depth 32, 200 nodes, 400 edges, 50 000 values. A local file over 1 MiB is never read: `<name>: <size> bytes exceed the limit of 1048576 (1 MiB); not read`.

Refusals to expect

- Gate codes and text: `EMPTY` *paste a Workflow IR document or choose a local JSON file*; `NOT_JSON` *not JSON: ... YAML, code and archives are never accepted*; `NOT_AN_OBJECT`; `T00_DEEP` *nesting deeper than 32 levels*; `T00_MANY_NODES` *<n> nodes exceed the limit of 200*; `T00_MANY_EDGES`; `CREDENTIAL_SHAPED`; `REMOTE_REFERENCE`.

- Draft refusals, rendered as `code · reason: draft_refused` for a document that does not validate under the composer's adapter, for a contract version the server does not support, for a declared version disagreeing with the one named, or for typed input it rejects; `draft_duplicate` for the same document kept twice; `supersession_refused` when the named predecessor is not the head of its lineage; `registration_link_refused` when the digest is not that of the registration named, because a link binds one exact record.
- With no drafts file configured, all three draft operations answer the typed gap `workflow_drafts` rather than an error, and the controls say so.
- Server 422 `workflow_too_complex`, e.g. *workflow: nodes 201 exceeds the limit 200*, *workflow: depth 33 exceeds the limit 32*; a raw body over 1 MiB is a 413 before parsing; an unsupported version on Adapt is 422 `unsupported_contract_version`. Rendered as *The server refused: <code> · <message>*.

Source: `features/bring/BringYourWorkflowScreen.tsx:120-888`; `features/bring/gate.ts:11-236`; `backend/.../api/workflows.py:31-131`; `workflow_limits.py:35-102`; drafts: `services/studio_v2.py:1647-1845` and `UGENCE_STUDIO_WORKFLOW_DRAFTS_PATH` (`config.py:98-101`); `frontend/tests/fixtures/bring.validate.json`, `bring.compare.json`.

What is not a screen, and why

- **Granting, revoking:** implemented on the worker behind its identity gate (AW-2 to AW-5) and not served until the identity adapter is validated end to end against a real enterprise issuer (AP-3 controlling, §18). The screens built for them under AW-1 were withdrawn with that ruling's reversal.
- **Issuing, activating:** the authority plane's other two writes, on stores the worker does not compose (AW-2), behind AP-3.
- **Authorizing, clearing, executing:** runtime authority, with ActionGate, the Autonomous Control Plane and the runtime. Refused on every front end by SD-2 and AP-4.
- **Module composition:** providers, hooks and seam files are environment variables read before the port binds. Shown read-only on Status; set nowhere in a browser.
- **The console's module registry:** behind a withheld route; struck from the studio by MS-1.
- **Risk Authority:** its own package records and one administrative emergency stop, with no web surface; the studio is forbidden to import it.
- **A general agent uploader or framework converter:** refused by BW-1. Bring Your Workflow accepts Ugence Workflow IR only. Conversion is an offline command-line tool with no screen (`packages/tooling/workflow-converters` , ruling CV-1 to CV-5, §23): the n8n and BPMN 2.0 converters emit a DRAFT policy pack, a conversion report and a preview Workflow IR labelled `PREVIEW_UNAPPROVED` that this screen can inspect; LangGraph, CrewAI and AutoGen stay deferred until a declarative export exists.