

Ugence + ServiceNow

Independent Action-Level Authority for Governed Agentic Workflows

Extending ServiceNow-governed workflows with independently verifiable authorization, operational clearance, execution reconciliation, and governed-value evidence.

 ServiceNow

 Ugence decision / authority

 Ugence execution / assurance

 External target

 Human review

 Hold / Block / Escalate

Rakesh Mohan · Founder, Ugence Labs · ugence.ai · August 2026

All ServiceNow integrations are PROPOSED — no connector ships today. All scenarios are illustrative, not customer deployments.

Two levels for two audiences — on every major scenario

Level 1 — Layman view

For business, operations and partnership stakeholders. Each scenario first answers, in plain language:

- The enterprise problem and why it matters
- What could go wrong when an AI agent acts
- What the AI agent wants to do
- Where control is required
- What ServiceNow already provides
- What Ugence proposes to contribute
- What the enterprise receives at the end

Level 2 — Technical architecture view

For solution architects. Each scenario then shows:

- What business data enters, and where it originates
- What each module receives or references
- What it checks or decides
- The separate governance artifact it emits
- What information passes over every arrow
- What executes the action
- What operational evidence returns
- How the observed effect is reconciled
- What is recorded or referenced in ServiceNow

How to use this briefing

This is a detailed landscape presentation designed for screen-based discussion and technical reference. The meeting walkthrough can focus on the executive framing, UC-5 autonomous change execution, the pilot proposal, and the discovery questions; the remaining sections provide supporting architecture and use-case detail.

How this briefing is organized — navigate by slide number

I Executive framing

II Data journey

III Module responsibilities

IV Lead UC-5 walkthrough

V Additional scenarios

VI Full use-case portfolio

VII Honest overlap & differentiation

VIII Enterprise Governed Value

IX Pilot proposal

X Discovery & next step

I

Executive framing

The problem, the discovery questions, what ServiceNow provides, and the partnership boundary

When an AI agent acts, can you prove it stayed in bounds?

“When an AI agent takes a consequential action, how can the enterprise prove that this exact action was authorized, remained safe at execution time, and produced only the intended effect?”

1. Wrong target

The agent acts on the wrong system, resource, or record — a valid-looking action pointed at the wrong place.

2. Conditions change

Time passes between approval and execution; a freeze begins, a dependency degrades, cost limits shift.

3. Effect ≠ record

What actually happened in the environment does not match what the record says was authorized.

A simple business example. An AI operations agent believes the online-checkout service needs more capacity. Scaling may prevent an outage — but scaling the wrong cluster, on stale information, during a change freeze, or past the approved cost ceiling could raise cost or disrupt another service. The enterprise wants the speed of autonomy AND proof that the change that ran is the one that was approved, for the target approved, at a moment that was safe.

Why an ordinary approval may not be enough — as questions, not claims

?

Was the approval for this exact target and payload — or for a class of action?

?

Is the approval still valid at the moment of execution?

?

Did the proposed action or payload change after approval?

?

Are live operational conditions still safe right now?

?

Did execution affect only the intended resource?

?

Can the enterprise prove the complete chain afterward?

What ServiceNow already provides — the proposal is complementary

AI Control Tower

Discovers, governs, secures and measures AI across the enterprise; real-time enforcement.

Action Fabric

A governed system of action (MCP / A2A) — every action identity-verified, permission-scoped, auditable.

Agent Deviation Detection

Flags when an agent strays from its authorized role at runtime.

AI Risk & Compliance

Multi-framework control mapping (EU AI Act, NIST AI RMF and others).

Change Management + CMDB

Approvals, change/blackout windows, conflict detection, business context.

NVIDIA OpenShell

Central AICT policy enforced at runtime on file, command and network access.

Adoption, impact & ROI

AI adoption, business impact, realized value and ROI measurement.

Approvals & workflow

Enterprise approval and workflow mechanisms across the platform.

The partnership boundary — who contributes what

ServiceNow remains responsible for

- Enterprise system of record
- Workflow and case management
- CMDB and business context
- Enterprise approvals
- Platform governance
- Action Fabric
- Platform execution pathways
- AI Control Tower monitoring & controls
- Value and ROI reporting

Ugence proposes to contribute

PROPOSED

- Binding decision artifacts
- Trusted-evidence linkage
- Exact-action and target authorization
- Time- and scope-bounded authority
- Independent operational clearance
- Governed execution coordination
- Execution receipts
- Observed-effect reconciliation
- Authority lifecycle & reassessment
- Cross-stage evidence lineage
- Developing governed-value attribution

II

Understanding the data journey

Four kinds of data, and the same journey told first in plain language, then technically

Four kinds of data — kept distinct throughout

1 · Original business data

Change, incident, CI, employee, entitlement, control evidence, amount, target system, requested action

ServiceNow (system of record)

2 · Derived governance artifacts

Decision record, model / action / risk authorization, clearance verdict, execution receipt, assurance result, revocation / reassessment signal

Ugence — each a SEPARATE artifact

3 · Observed operational data

Infrastructure state, service health, blackout window, dependency condition, account status, execution result, post-action effect

ServiceNow + operational telemetry


4 · ServiceNow records

Remain the system-of-record context; receive statuses, receipt references and summarized outcomes

ServiceNow (via PROPOSED integration)

“Each Ugence module reads or references the necessary business context and emits a separate governance artifact. It does not silently rewrite the originating ServiceNow record.”


End-to-end, in plain language



Example information added

- Proposed: "scale checkout 12 → 18"
- Decision: approved under delegated scaling authority
- Permission: this cluster, service, window
- Clearance: no freeze, deps healthy, cost within limit
- Receipt: request accepted, completed
- Effect: 18 active, nothing else changed
- Result: execution matched authorization

The technical canonical workflow — modules that may participate



Each layer speaks a distinct decision verb; authority never leaks across layers. Not every module appears in every scenario.

III

Module responsibilities

A business-readable glossary, and the repository-verified architecture facts

Architecture facts we hold to (verified against the source of record)



Decision Authority owns the binding decision

It also owns execution and reconciliation records, and the context-envelope records (CER). AI is barred as an authorizing principal.



Agent Runtime coordinates governed execution

It owns canonical execution state and invokes domain executors / providers / tools. It is not the "CER," and it creates no authority.



Cloud Scaling Operations is a domain executor

It is invoked by Agent Runtime and returns its result to Agent Runtime — it is not a preceding authorization gate.



RA-8 compares intent, evidence and effect

It checks authorized intent against the execution receipt and observed effect. It never retroactively authorizes an action.



Authority stays scoped and revocable

RA-6 is the sole writer of authority lifecycle (revoke / supersede / expire). Enforcement stays read-only.



Fail-closed at every boundary

Missing approval, evidence, clearance or execution confirmation is never treated as permission.

Module glossary (1 of 3) — policy, evidence, decision, signed authority

Module	Layman responsibility	Receives / references	Checks / decides	Emits	Does not do
Policy Workflow Compiler	Turns approved policy into deterministic, checkable constraints	Approved, structured policy pack (compile-time, offline)	Compiles policy to constraints	Digest-addressed governed-workflow artifact	No binding decision; authorizes / clears / runs nothing
RA-5 Trusted Evidence Admission	Ensures a control is met by trusted, re-checked evidence	Control evidence references	Is each control satisfied by trusted evidence? (a caller “pass” is inert)	Evidence-derived, re-checked control result	Add a second authority signature
Decision Authority	Confirms a delegated authority — not the AI — approved this class of decision	Proposed action, delegated authority, constraints; owns CER	Is there a binding decision within scope?	A separate, immutable decision record	Execute; inspect live conditions; let AI self-authorize
Risk Authority	Converts an approved risk decision into signed, time-limited, scoped permission	An allow-family risk decision	Whether to mint authority; scope $\leq$ decision scope	A signed, scoped, time-limited authorization (tamper-evident)	Execute the action

Module glossary (2 of 3) — model, exact action, sequence, clearance

Module	Layman responsibility	Receives / references	Checks / decides	Emits	Does not do
Model Authority	Decides which model may handle this specific request, now	Request context, approved model policy	Per-request eligibility	ALLOW / DENY / HOLD / ESCALATE, with governed fallback and expiry	Execute the request; replace platform provider approval
ActionGate	Converts an approval into permission for one exact action on one exact target	Decision record, target identity, proposed action	Same authorized action? still valid?	AUTHORIZED / DENIED / INDETERMINATE (uncertainty → INDETERMINATE)	Judge operational safety; own execution
StoryGraph	Flags when individually-harmless steps add up to a harmful capability	A multi-step plan across steps	Do benign steps assemble a harmful capability?	OBSERVE / ESCALATE (advisory)	Authorize or execute
Action Clearance (ACP)	The final “is it safe right now?” check	Already-authorized action + current conditions	Blackout, conflict, unhealthy dependency, operational hold	CLEAR / HOLD / BLOCK / ESCALATE (subtractive)	Create or broaden authority; dispatch execution

Module glossary (3 of 3) — coordination, actuation, lifecycle, assurance

Module	Layman responsibility	Receives / references	Checks / decides	Emits	Does not do
Agent Runtime	Coordinates the governed execution and carries its state	A governed request; invokes providers / domain executors	Execution lifecycle (retry / timeout / recovery)	Execution receipt; owns canonical execution state	Create authority; author policy; authorize; mint clearance
Cloud Scaling Operations	Carries out a bounded infrastructure change under strict controls	A scaling instruction + externally minted execution authorization	Readiness; bounded execution (dry-run by default)	Execution outcome + audit, returned to Agent Runtime	Mint its own authority; act as a preceding gate
RA-6 Authority Lifecycle	Keeps authority current — revokes, supersedes, expires	A reassessment signal when conditions change	Whether to revoke / supersede / expire	A lifecycle mutation (bites at the next pre-effect recheck)	Execute or authorize
RA-7 Trajectory Assurance	Watches an in-flight execution for drift	The in-flight execution (neutral observation)	Is the trajectory drifting?	NORMAL / ESCALATED / UNKNOWN → reassessment signal	Mint authority
RA-8 Execution Assurance	Checks whether reality matched what was authorized	Authorized action, execution receipt, observed state	Did only the intended target change, within scope?	matched / mismatch / partial / unknown	Retroactively legitimize an unauthorized action

IV

Lead scenario — autonomous change execution

UC-5: the full two-level walkthrough, end to end

The problem, in plain terms

Illustrative scenario. *Change CHG0048217 · online checkout · production Kubernetes cluster · scale the checkout service from 12 to 18 instances · sustained demand and latency pressure · constraints: approved cost ceiling, permitted change window, healthy dependencies, rollback capability. This shows how the proposed integration could operate — it is not a claimed Ugence customer deployment.*

Why the AI proposes scaling

Checkout latency is rising under sustained demand; more capacity may prevent an outage.

What is at risk

Online checkout — a revenue-critical business service — and any service that shares the cluster.

What could go wrong

Scaling the wrong cluster or service, or acting on stale information, could disrupt another service or exceed cost.

Why a valid approval may still be unsafe

A freeze may have begun, a dependency may be unhealthy, or projected cost may now exceed the ceiling.

Why the effect must be checked

The record must reflect what actually happened — only the intended instances changed, nothing else.

What the enterprise wants


Autonomy with proof: the exact approved change, at a safe moment, with a verifiable, reconciled result.

What business data enters the workflow


Data category	Example	Authoritative source
Business record	Change request CHG0048217 and its approval state	ServiceNow Change Management
Target context	Cluster, checkout service, requested capacity 12 → 18	CMDB / ITOM / cloud platform
Live conditions	Freeze window, dependency health, service health	ServiceNow + operational telemetry
Governance context	Delegated scaling authority, cost / risk limits, expiry	Approved enterprise policy
Intended outcome	Lower latency without exceeding cost or risk limits	Business objective

Business-readable fields only — not internal schemas. Original business data (rows 1–2) is distinct from live conditions (row 3) and from any governance artifact Ugence later emits.

The business journey (plain language)



Technical module workflow — Agent Runtime invokes the executor



What stops the action — and what happens then

Missing evidence / approval

The workflow does not proceed; uncertainty is not treated as permission.

Changed target or payload

ActionGate returns DENIED / INDETERMINATE — the authorization no longer matches.

Unsafe live conditions

Action Clearance returns HOLD / BLOCK / ESCALATE on freeze, conflict or cost.

Authority revoked / expired

RA-6 revocation bites at the next pre-effect recheck; further action is blocked.

Execution mismatch

RA-8 flags mismatch / partial / uncertain — the effect did not match the authorization.

Governance not wired

Agent Runtime fails closed — consequential transitions are blocked by default.

Fail-closed: if a mandatory approval, evidence item, live-safety check or execution confirmation is missing, the workflow does not treat uncertainty as permission.

Module-by-module walkthrough (1 of 2)

Module	Receives	Question it answers	Adds	Sends next	If information is missing
Decision Authority	Proposed change, delegated authority, constraints	Is there a binding decision within delegated scope, by a non-AI principal?	A binding decision record	To ActionGate	No delegated authority → no binding decision; stop
ActionGate	Decision record, cluster identity, proposed action	Is this the exact authorized action, still valid?	An exact-target action authorization	To Action Clearance	Mismatch or uncertainty → INDETERMINATE; stop
Action Clearance	The authorization + current live conditions	Is it operationally safe right now?	A CLEAR / HOLD / BLOCK / ESCALATE verdict	To Agent Runtime (on CLEAR)	Unsafe or unknown live state → HOLD / BLOCK

Module-by-module walkthrough (2 of 2)

Module	Receives	Question it answers	Adds	Sends next	If information is missing
Agent Runtime	The governed request (on CLEAR)	How is execution coordinated safely?	An execution receipt; owns canonical execution state	Invokes Cloud Scaling Operations; carries receipt to RA-8	No governance wired → fails closed; blocks
Cloud Scaling Operations	A scaling instruction + external execution authorization	Is it ready, and bounded to the authorization?	Execution outcome + audit	Returns result to Agent Runtime	No execution authorization → dry-run only; no live change
RA-8 Execution Assurance	Authorized action, receipt, observed cluster state	Did only the authorized instances change, within scope?	matched / mismatch / partial / unknown	Status + receipt refs to ServiceNow	Cannot verify → uncertain / manual review; never auto-authorize

Business outcome, and the Cloud Scaling maturity picture

Business & pilot outcome

- Benefit: latency relief without waiting on a human
- Risk prevented: wrong cluster, freeze-time change, cost overrun
- Evidence produced: decision · authorization · clearance · receipt · effect-match — one lineage per change
- ServiceNow retains: change record, CMDB, windows, execution pathway
- Ugence contributes: exact-action authority, independent clearance, effect reconciliation
- Pilot metrics: change success rate; rollback frequency; unauthorized-target attempts blocked; effect-match rate

Possible governed-value measures

UNDER DEVELOPMENT

A developing, cross-cutting capability (not an authorization gate):

- Infrastructure cost held within the ceiling
- Service availability preserved
- Latency improvement attributed to the governed action
- No change-window, risk or availability violation

ServiceNow AI Control Tower already measures adoption, business impact, realized value and ROI.

Cloud Scaling maturity — four separate dimensions (do not conflate)

Core Cloud Scaling Controller

IMPLEMENTED

Production validation

PILOT PENDING

Additional agentic-AI capabilities

UNDER DEVELOPMENT

ServiceNow integration

PROPOSED INTEGRATION

V

Additional enterprise scenarios

UC-11 · UC-6 · UC-3 · UC-4 — each at both levels

The problem and the data entering the workflow

An AI agent proposes to deploy an emergency patch to production servers. Speed matters — but patching the wrong set of servers, patching a business-critical system at the wrong moment, or continuing after the risk picture worsens mid-rollout can cause the very outage the patch was meant to prevent.


Illustrative: *VUL0007731 → remediation task → change record · critical RCE on a web tier · 40 CIs in “payments-web” · emergency patch · staged rollout · ability to stop mid-flight.*

Data category	Example	Source
Business record	Remediation task + change record, approval state	ServiceNow Vulnerability Response + Change
Target context	40 CIs in payments-web, the patch	CMDB / Security Operations
Live conditions	Business-criticality, maintenance window, risk posture	ServiceNow + telemetry
Governance context	Delegated emergency-patch authority, expiry, revocability	Approved enterprise policy

Why staged rollout matters

- Each stage is cleared for live conditions and business-criticality before it runs
- If the risk posture worsens, authority can be revoked before the next server is touched
- Only the authorized configuration items are ever changed

Technical module workflow — with mid-rollout revocation



Revocation is bounded-latency: it stops the next stage at the pre-effect recheck, not one already in progress. Illustrative example.

Human-control boundary, and business & pilot metrics

Human-control boundary

- Autonomous: per-stage clearance and patching INSIDE the authorized CI set and window
- Forces HOLD / ESCALATE: peak-trading or criticality hold, window closed, expired authority, a revocation signal
- Human-binding: the emergency-patch authority delegation; any CI outside the authorized set
- Limits: named CI set, staged rollout, reversibility, revocability mid-flight
- Fail-closed: missing approval, evidence, clearance or execution confirmation is never permission

Business outcome & pilot metrics

Benefit: faster mean-time-to-remediation for critical vulnerabilities.
Risk prevented: out-of-scope or business-critical systems patched at the wrong moment; runaway rollout after risk worsens.
Evidence: per-stage authorization, clearance, receipts, revocation record, effect-match.

Possible pilot metrics:

- Mean-time-to-remediation
- Unauthorized-target attempts blocked
- % of stages correctly stopped on revocation
- Effect-match rate


Governed-value (developing): exposure-window reduction attributed to governed remediation, with preserved availability.

The problem, and the business journey


An employee asks a self-service assistant for elevated access — for example, the ability to both create and approve payments. An AI agent could fulfil it in seconds. But granting an entitlement that breaks segregation of duties, or that the requester’s current risk posture should block, can create fraud exposure that is hard to unwind.

Illustrative: RITM0102934 · a finance analyst who already holds “create payment” requests “approve payment” in the ERP · an SoD rule prohibits one person holding both.

ServiceNow already provides Service Catalog request fulfilment, entitlement management, and Veza-based identity and least-privilege controls. Ugence proposes exact-entitlement authorization and an independent SoD clearance.



Technical module workflow — segregation-of-duties veto



The decisive control is the SoD conflict at clearance time — the action can be perfectly authorized and still be blocked. Illustrative.

Human-control boundary, and business & pilot metrics

Human-control boundary

- Autonomous: fulfilling entitlements with NO SoD conflict and a clear risk posture
- Forces HOLD / ESCALATE / BLOCK: any SoD conflict, elevated risk posture, disabled account, expired authority
- Human-binding: the access-grant authority delegation; any entitlement outside scope
- Limits: exactly one entitlement per authorization; reversibility; SoD rule as a hard veto
- Fail-closed: if SoD state or approval cannot be established, the request is not treated as permitted

Business outcome & pilot metrics

Benefit: fast self-service access without a queue.

Risk prevented: SoD violations and fraud exposure; silent AI self-grant.

Evidence: decision, exact-entitlement authorization, SoD clearance, effect-match (or a clean “no change” on block).

Possible pilot metrics:

- Access-policy violations prevented
- % of grants with matched effect
- Time saved per request
- Unauthorized-entitlement attempts blocked

Governed-value (developing): fraud-exposure reduction attributed to blocked SoD conflicts, with preserved compliance.

A concrete, regulated illustrative example

Illustrative (privileged system access). An AI use case classified high-risk is about to grant privileged production access automatically. The policy is approved and controls are on file — but at the moment of action, is every required control currently satisfied by trusted, re-checked evidence, or is the system relying on a stale, self-asserted “pass”? Acting on out-of-date compliance evidence is exactly what audits punish.

Discovery questions (not claims about ServiceNow)

- Is each control satisfied by trusted, re-checked evidence — not a caller-asserted pass?
- Is the resulting permission a signed, scoped, time-limited artifact a downstream gate can verify?
- Did the action that ran stay within the authorized scope?


System-of-record stays with ServiceNow

AI Control Tower and AI Risk & Compliance remain the governance systems of record — classification, multi-framework control mapping (EU AI Act, NIST AI RMF and others), and enforcement.

Ugence proposes an independently verifiable, action-level authority artifact linked to CURRENT evidence — so a stale or self-asserted control cannot wave an action through.

Anything not established in ServiceNow documentation is treated as a discovery hypothesis, not a claim.

Technical module workflow — evidence-linked signed authority



Positioning, control boundary and metrics

Positioning (honest)

- ServiceNow AI Control Tower & AI Risk and Compliance remain the governance systems of record — this is a real, shipped strength.
- Ugence proposes a signed, scoped, time-limited, action-level authority whose validity depends on trusted, currently-satisfied evidence.
- Whether ServiceNow already emits an equivalent independently verifiable, evidence-fresh per-action artifact is a discovery hypothesis to confirm with architects — not an assumed gap.
- Human-binding: the policy approval and the authority delegation. Fail-closed: stale or self-asserted evidence is never a satisfied control.

Outcome & pilot metrics

Benefit: high-risk automation that stays inside an evidence-backed, approved boundary.

Risk prevented: acting on stale compliance evidence; unbounded high-risk actions.
Evidence: re-checked control result, signed authorization, exact-action authorization, effect-match — an audit-ready chain.

Possible pilot metrics:

- % of high-risk actions with fresh-evidence backing
- Stale-evidence blocks
- Effect-match rate

Governed-value (developing): audit-finding reduction attributed to evidence-fresh enforcement, with preserved compliance constraints.

The problem, and the business journey

Enterprises now let external AI agents (for example Claude or Copilot) take real actions on enterprise systems. ServiceNow Action Fabric opens its system of action to these agents and routes every action through AI Control Tower; with NVIDIA OpenShell, policy is enforced at runtime on file, command and network access.

The remaining question is narrow: for each business action an external agent takes, is there an independent, verifiable record that THIS exact action was authorized — and does anything notice when individually-allowed steps add up to something that should not be allowed?

Ownership (unchanged). Action Fabric owns the governed ServiceNow action pathway; NVIDIA OpenShell owns its runtime / sandbox enforcement. Ugence contributes authority, exact-action binding, evidence lineage, clearance, sequence-risk and assurance artifacts — composition, not replacement.

External AI agent proposes a step

via Action Fabric

Allowed model & exact action?

bound to payload

Do the steps together look harmful?


sequence risk

Enforced by the platform

Action Fabric + OpenShell

Independent record to ServiceNow

Technical module workflow — composition, not ownership



Positioning, ownership and metrics

Ownership & positioning

- Action Fabric owns the governed ServiceNow action pathway; every action is identity-verified, permission-scoped and auditable.
- NVIDIA OpenShell owns runtime / sandbox enforcement on file, command and network access.
- Ugence contributes: an independent, signed, exact-payload authorization; sequence-risk advisory across the plan; trajectory assurance — composition and partnership, not replacement.
- Ugence does NOT claim to own Action Fabric or OpenShell execution, nor unique cross-runtime enforcement.

Outcome & pilot metrics

Benefit: external agents get real work done under a verifiable, independent authorization trail.

Risk prevented: an unauthorized payload, or a harmful combination of individually-allowed steps, moving value unnoticed.

Evidence: model authorization, exact-payload authorization, sequence-risk signal, trajectory assurance.

Possible pilot metrics:

- Unauthorized-payload attempts blocked
- Sequence escalations caught
- % of actions with an independent authorization record

Governed-value (developing): prevented-loss attributed to sequence and payload controls, with preserved service constraints.

VI

The full use-case portfolio

All 12 scenarios — grounded in ServiceNow products, with proposed Ugence governance extensions

Pilot-ready operational scenarios

Group A — pilot-ready operational

UC-5 · Autonomous change execution

Problem: Scale/restart/config changes made autonomously

ServiceNow anchor: Change Management + ITOM

Ugence proposes: Exact-CI authority, independent clearance, effect reconciliation

Autonomy: Autonomous within window & cost; freeze/conflict force stop

Discovery: Which change + CMDB events are authoritative?

PROPOSED INTEGRATION

UC-6 · Access provisioning with SoD

Problem: Elevated access fulfilled by an assistant

ServiceNow anchor: Service Catalog / Employee Center + Veza

Ugence proposes: Exact-entitlement authority + independent SoD clearance

Autonomy: Autonomous when no SoD conflict; conflict → block

Discovery: Which entitlement class suits a bounded pilot?

PROPOSED INTEGRATION

UC-11 · Vulnerability remediation

Problem: Emergency patch to production CIs

ServiceNow anchor: Security Operations — Vulnerability Response

Ugence proposes: Exact-CI binding, per-stage clearance, mid-rollout revocation

Autonomy: Per-stage autonomy; revocation stops next stage

Discovery: Which critical-vuln workflow for a pilot?

PROPOSED INTEGRATION

Practical scenarios requiring bounded autonomy

Group B — bounded autonomy

UC-7 · Refunds & credits

Problem: AI issues a refund or credit

ServiceNow anchor: Customer Service Management (Now Assist)

Ugence proposes: Amount+account binding; reconcile executed = authorized

Autonomy: Within a delegated dollar threshold

Discovery: Which refund band is safe to pilot?

PROPOSED INTEGRATION

UC-8 · Procurement / PO issuance

Problem: AI places a purchase order

ServiceNow anchor: Sourcing & Procurement Operations

Ugence proposes: Compiled policy, supplier-evidence check, exact PO binding

Autonomy: Threshold-bound; above → human authority

Discovery: Which spend threshold and category?

PROPOSED INTEGRATION

UC-10 · Agentic hiring (human-binding)

Problem: AI assists screening & scheduling

ServiceNow anchor: HR Service Delivery / Recruitment

Ugence proposes: Keep hire/reject human; immutable decision record

Autonomy: Assistive only; decision stays human-binding

Discovery: Which step may be automated vs human-bound?

PROPOSED INTEGRATION

Strategic governance scenarios

Group C — strategic governance

UC-2 · Per-request model authorization

Problem: Regulated data about to reach a model

ServiceNow anchor: AI Control Tower model/provider governance

Ugence proposes: Per-request ALLOW/DENY/HOLD/ESCALATE + fallback/expiry

Autonomy: Per-request binding, not a static allowlist

Discovery: Which data classes trigger per-request control?

PROPOSED INTEGRATION

UC-3 · High-risk action enforcement

Problem: High-risk AI action on stale evidence

ServiceNow anchor: AI Control Tower — AI Risk & Compliance

Ugence proposes: Signed, scoped, evidence-fresh action authority

Autonomy: Only while trusted evidence is current

Discovery: Which high-risk use case for a pilot?

PROPOSED INTEGRATION

UC-4 · External-agent governance

Problem: External agents act on the system of action

ServiceNow anchor: Action Fabric + AI Agent Fabric + OpenShell

Ugence proposes: Independent exact-payload authority, sequence-risk, assurance

Autonomy: Composes with platform enforcement

Discovery: Which external-agent workflow to pilot?

PROPOSED INTEGRATION

UC-12 · Data-boundary governance

Problem: Agents pull enterprise data into context

ServiceNow anchor: Workflow Data Fabric + AICT privacy

Ugence proposes: Minimum-necessary context; token accounting; fail-closed

Autonomy: Governs exactly what crosses the model boundary

Discovery: Which data domains need boundary control?

PROPOSED INTEGRATION

Emerging or future scenarios

Group D — emerging / future

UC-1 · Autonomous security containment

Problem: Autonomous host isolation / account disable / IP block

ServiceNow anchor: Security Incident Response — Tier 2 SOC AI Specialist

Ugence proposes: Exact-target authority, live-safety clearance, effect reconciliation

Autonomy: Forward-looking; not shipped

Discovery: When the specialist ships, where does authority bind?

ANNOUNCED / FUTURE

UC-9 · Governed multi-agent workforce

Problem: A team of agents runs an end-to-end process

ServiceNow anchor: AI Agent Orchestrator / Autonomous Workforce

Ugence proposes: Least-privilege team plan granting nothing; per-action authorization + sequence risk

Autonomy: Composition grants no authority

Discovery: Which multi-agent process to bound first?

PROPOSED INTEGRATION

VII

Honest overlap & differentiation

Name the overlaps first — then the narrow, artifact-level differentiation

Confirmed overlap zones — discussed as overlaps, not gaps

Zone	ServiceNow capability	Ugence proposed edge (narrow)
ActionGate ↔ Action Fabric / AICT	Governed, identity-verified, auditable action execution	Per-payload/target authorization re-checked at commit (discovery hypothesis)
RA-7 ↔ Agent Deviation Detection	Flags when an agent strays from its role	Independently verifiable trajectory signal into an authority lifecycle
RA-6 ↔ real-time shutdown	Can stop a misbehaving agent	Authority-lifecycle revoke/epoch; enforcement stays read-only
Model Authority ↔ Skill-Kit approval	Restricts which models run	Per-request binding decision with governed fallback + expiry
AICT + NVIDIA OpenShell	Central policy enforced at runtime across file/command/network	Ugence does NOT claim unique cross-runtime enforcement — differentiation is the artifact

Absence from public documentation is never treated as proof ServiceNow lacks a property — those points are questions for ServiceNow architects.

The narrow, defensible edge — properties of one authority artifact

Ugence does not claim a governance gap or unique cross-runtime enforcement. The differentiation is confined to the properties of an independent, action-level authority-and-evidence artifact:

1 **Independently verifiable**

by a party other than the executor

2 **Bound to the exact business action & target**

not a class of action

3 **Linked to trusted evidence**

stale or self-asserted evidence is inert

4 **Scoped & time-limited**

scope never exceeds the decision

5 **Re-checked before consequential execution**

at commit time

6 **Separated from live operational clearance**

independent, subtractive veto

7 **Connected to execution receipts**

what actually ran

8 **Reconciled against observed effect**

intent vs reality

9 **Usable in a cross-stage evidence lineage**

one story per action

VIII

Enterprise Governed Value

A developing, cross-cutting capability — evidence-backed outcome attribution

Developing & cross-cutting — not an authorization gate

UNDER DEVELOPMENT

ServiceNow already measures value. AI Control Tower measures AI adoption, business impact, realized value and ROI. This proposal does not duplicate or replace that.

Ugence proposes attribution. Evidence-backed outcome attribution that connects approved objectives, governed-execution evidence, attributable cost, observed outcomes, and preserved risk, compliance, quality and service constraints.

Concrete example — UC-5 checkout scaling

Approved objective Reduce checkout latency	Baseline Current latency, capacity, cost	Governed action Scale 12 → 18 instances	Execution evidence Authorization · clearance · receipt
Observed outcome Latency improved; service healthy	Attributable cost Additional infrastructure expense	Preserved constraints No window / risk / availability violation	Governed-value question Enough attributable value without unacceptable trade-offs?

Clearly under development. It reuses the governance evidence the pipelines already emit — it adds attribution, not a new gate.

IX

Pilot proposal

One bounded technical pilot around UC-5 — no broad production autonomy

A bounded progression — earn autonomy step by step



Boundary and measurements

Possible pilot boundary

- One controlled scaling workflow (UC-5)
- A limited set of noncritical or carefully selected CIs / clusters
- Restricted action types
- Explicit authority and cost limits
- Defined rollback requirements
- No broad production autonomy at pilot start

The pilot proves the authorization → clearance → execution → reconciliation chain on a small, reversible footprint before any scope increase.

Possible measurements

- % of actions bound to the exact approved target
- Authorization mismatches blocked
- Unsafe live conditions detected
- HOLD / BLOCK / ESCALATE accuracy
- Execution-to-effect match rate
- Change success rate
- Rollback rate
- Service availability
- Time saved
- Infrastructure cost
- Attributable governed value (developing)

X

Discovery & next step

Constructive questions, and a request for one joint technical session

Constructive questions to work through together

1 Which ServiceNow records are authoritative for each scenario?

2 What events expose the approval, target and execution context?

3 Where could an independent action-level authorization be evaluated?

4 Where could a live operational clearance verdict be consumed?

5 How should execution receipts and effect results be associated with the originating record?

6 Which controls already provide equivalent granularity?

7 Which use case is most suitable for a bounded pilot?

8 What should remain human-binding?

9 What data may leave the ServiceNow boundary?

10 What evidence would be required for production acceptance?

THE NEXT STEP

One joint technical discovery session

to select and define one bounded pilot — most likely UC-5, autonomous change execution.

We are not asking for

a broad partnership commitment before technical fit is established.

We are asking for

a working session with your architects to map records, events and integration points, and to agree a small, reversible pilot.

All ServiceNow integrations PROPOSED · all scenarios illustrative · composition, not replacement.

Rakesh Mohan · Founder, Ugence Labs · ugence.ai

Maturity & availability legend, and standing disclaimers

IMPLEMENTED	Code exists and is verified at interface level.
REFERENCE-GRADE	Ships reference implementations; not production-validated.
PILOT PENDING	Awaiting production / pilot validation.
UNDER DEVELOPMENT	Additional capabilities actively being built.
PROPOSED INTEGRATION	A ServiceNow integration adapter is design intent; no connector ships.
ANNOUNCED / FUTURE	A named future capability; not yet shipped.
DESIGN-ONLY	Architecture / design; not a package.

Standing disclaimers. All ServiceNow integrations are PROPOSED — no connector ships. All scenarios are illustrative, not customer deployments. Maturity dimensions are not conflated: a multi-module workflow carries per-module maturity. Unverified differentiation is a discovery hypothesis, never a claim that ServiceNow lacks a capability.

Source hierarchy & citations

1 · ServiceNow Docs	Existing product behavior and workflow mechanics (docs.servicenow.com)
2 · ServiceNow release notes	Versions and availability (what ships when)
3 · ServiceNow Newsroom	Announcements, partnerships and future availability
4 · ServiceNow Community	Supporting or explanatory material
5 · NVIDIA documentation	OpenShell technical behavior (seccomp + Landlock LSM + network namespaces; not eBPF)
6 · Third-party press	Corroboration of announcement facts only (e.g. Dec-2026 SOC availability)

Key anchors: AI Control Tower · Action Fabric · AI Agent Fabric · AI Risk & Compliance · Change Management · Security Incident Response · Vulnerability Response · Service Catalog · HRSD · Sourcing & Procurement · NVIDIA OpenShell.

Ugence module behavior is grounded in the source-of-record documents (catalog v1.2 and the walkthrough companion) at the guarantee level; no internal repository detail is exposed here.